

Chapter 1

Introduction

1.1 From Hit – to – Lead: The Bottleneck in Modern Drug Discovery

Although the sequencing of the human genome and the advent of high-throughput screening (HTS) promised a flood of new therapeutics, the average cost of bringing a small-molecule drug to market has risen to \approx \$2.6 billion and now spans 10–15 years **Jayatunga2024**. The *hit-identification* phase is particularly wasteful: historically only 0.1–0.3 % of compounds tested *in vitro* show any measurable activity against an intended protein target **BlancoGonzalez2023**. *Structure-based virtual screening* (SBVS) seeks to alleviate this attrition by replacing wet-lab assays with physics-based **molecular docking** simulations that predict the bound pose and free energy, ΔG_{bind} , of a ligand–receptor complex **Fan2019; Pagadala2017; Meng2011**.

Docking engines perform two coupled tasks: (i) an **exhaustive search** of the ligand’s translational, rotational and torsional degrees of freedom; and (ii) a **scoring function** that approximates ΔG_{bind} . Classical scoring functions decompose the free energy into additive physico-chemical terms

$$\Delta G_{\text{bind}} = w_{\text{vdW}}E_{\text{vdW}} + w_{\text{ele}}E_{\text{ele}} + w_{\text{solv}}E_{\text{solv}} + w_{\text{tor}}N_{\text{tor}} + C \quad (1.1)$$

where E_{vdW} and E_{ele} denote van-der-Waals and Coulombic contributions, E_{solv} the desolvation penalty, N_{tor} the torsional entropy term, and w_i empirically fitted weights **Dias2008; Stanzone2021**. Despite decades of refinement, Eq. (1.1) often overlooks receptor flexibility, solvent networks and entropic effects, yielding many false positives and negatives **Pagadala2017; Muhammed2022**.

1.2 Artificial Intelligence Meets Molecular Docking

The last decade has witnessed an explosion of **artificial intelligence** (AI) techniques—including machine learning (ML), deep learning (DL) and reinforcement learning (RL)—that augment or replace specific steps of the classical docking workflow **Gupta2021; Qureshi2023; Deng2022**. AI delivers value along three orthogonal axes:

- 1) **Surrogate Scoring:** Random forests (RF) **Sarkar2023**, gradient-boosted trees and graph neural networks (GNNs) **David2020** learn non-additive interaction patterns directly from crystallographic complexes, achieving root-mean-square errors below $1.5 \text{ kcal mol}^{-1}$ on the PDBbind core set—significantly outperforming physics-based Eq. (1.1) **Deng2022**.

- 2) **Intelligent Library Design:** RL agents sequentially assemble SMILES strings, using the ML score or docking free energy as a reward signal. This skews exploration towards synthetically viable, high-affinity areas of chemical space while adhering to medicinal chemistry principles. **Gupta2021; Bohr2020.**
- 3) **Active-Learning Screening:** Iterative cycles of docking, model retraining, and uncertainty-guided selection eliminate the necessity to simulate every molecule, reducing CPU time by up to 90% without sacrificing enrichment **Deng2022; BlancoGonzalez2023.**

1.3 Problem Statement

The rapid proliferation of publicly available chemical databases has greatly increased the search space for new medication possibilities. As of May 2025, the PUBCHEM repository—maintained by the NIH—hosts over 1.19×10^8 unique, canonicalized small molecules, each annotated with a standardized SMILES representation. The extensive chemical universe presents unparalleled prospects for virtual screening, while also imposing significant computational and methodological obstacles for conventional structure-based methods.

Molecular docking remains one of the most popular *in silico* techniques for hit identification, simulating ligand binding poses and estimating receptor–ligand affinity through approximated scoring functions **Fan2019; Pagadala2017; Muhammed2022.** However, docking a single ligand can consume several CPU-minutes, and scaling this operation across the entire PUBCHEM catalogue would require on the order of $\mathcal{O}(10^9)$ CPU-hours—rendering it infeasible for most academic and early-stage industrial research labs **Meng2011; Dias2008.**

Moreover, traditional scoring functions frequently exhibit flaws stemming from their failure to adequately account for receptor flexibility, solvation effects, and entropic contributions **Pagadala2017; Gao2017.** Consequently, a significant percentage of docked poses may produce false positives or negatives, exacerbating the inefficacy of brute-force virtual screening **David2020.**

To address these limitations, there is a critical need for **lightweight, scalable, and interpretable AI models** that can rapidly pre-filter large molecular libraries—ideally within hours on modest computational hardware. This pre-screening would favor the most promising candidates for high-fidelity simulations, such as docking or molecular dynamics (MD), therefore substantially decreasing both computing expenses and experimental burdens.

Recent research have suggested the incorporation of machine learning (ML) and deep learning (DL) as feasible replacements or enhancements to conventional docking procedures. Techniques such as surrogate scoring via random forests **Deng2022**, graph neural networks **Qureshi2023**, and reinforcement learning-based molecule design **Gupta2021; BlancoGonzalez2023** have demonstrated the potential to approximate docking scores or identify active compounds with comparable, and sometimes superior, accuracy at a fraction of the cost.

Consequently, the primary difficulty is not only how to replicate molecular interactions with greater precision, but how to do so *intelligently*—leveraging data-driven models to triage vast chemical spaces in a manner that is computationally tractable, chemically meaningful, and aligned with medicinal chemistry principles.

1.4 Proposed Multi-Stage AI Pipeline

Scaling virtual screening to the $\sim 10^8$ enumerated molecules now resident in PubChem demands an architecture that is *fast*, *chemically aware*, and *auditable*. Guided by best practice in structure-based drug design **Fan2019**; **Pagadala2017**; **Dias2008**, we devise a three-layer workflow (Fig. 3.1) that emulates the progressive filters medicinal chemists apply in hit-to-lead campaigns.

Layer 1: Descriptor-Based Random Forest.

Rationale. Classical "Lipinski-like" descriptors encapsulate the fundamental physicochemical parameters that regulate absorption, distribution, metabolism, and excretion (ADME) **Gao2017**. Their swift computability renders them optimal for an initial assessment of hundreds of millions of molecules.

Feature vector. Utilizing RDKit, we extract eleven scalar features: molecular weight (MW), logarithm of the partition coefficient (logP), topological polar surface area (TPSA), hydrogen bond donor and acceptor counts, heavy atom count, ring count, and the number of rotatable bonds, fraction sp^3 and BertzCT—then perform a corpus-wide z -normalisation,

$$x'_{ij} = \frac{x_{ij} - \mu_j}{\sigma_j}. \quad (1)$$

Model. A 200-tree Random Forest (max-depth 20, balanced class weights) gives a probability of activity

$$p_{\text{RF}}(x'_i) = \frac{1}{T} \sum_{t=1}^T \mathbb{I}[h_t(x'_i) = 1], \quad (2)$$

where $T = 200$ and h_t is the t^{th} tree. On a 32-core workstation the model screens $\sim 10^8$ molecules in <3 h. Decision-path inspection in individual trees provides transparent, regulator-friendly justifications **Tabassi2023**.

Layer 2: One-Dimensional Convolutional Neural Network (CNN).

Motivation. Non-linear coupling between descriptors—e.g. the dependence of passive permeability on both logP and TPSA—can confound purely tree-based methods **Fan2019**. A lightweight CNN adds universal approximation power without sacrificing throughput.

Architecture. The normalised vector is reshaped to 11×1 and passed through two Conv-ReLU-BatchNorm-Pool₂ blocks; the flattened output $z^{(2)}$ feeds a sigmoid-activated dense layer,

$$p_{\text{CNN}} = \sigma(W^\top z^{(2)} + b), \quad (3)$$

trained with Adam ($\eta_0 = 10^{-3}$), dropout 0.3 and early stopping. Cross-validation shows an average ROC-AUC gain of 1.5% over the RF, echoing observations that shallow CNNs can capture descriptor interactions missed by traditional QSAR **David2020; Deng2022**.

Layer 3: Reinforcement-Learning Molecule Generator (PPO).

Goal. After pruning $> 95\%$ of unequivocally poor candidates, we need fresh chemical matter that *optimises* predicted affinity yet remains synthesizable—precisely the setting where reinforcement learning excels **Gupta2021; BlancoGonzalez2023**.

Policy and reward. A Proximal Policy Optimisation agent autogenerates SMILES sequences; at each step t it receives

$$r_t = \underbrace{\sigma(p_{\text{RF}})}_{\text{activity}} - \lambda \underbrace{|\text{SA} - 5|}_{\text{synthetic access}} - \mu \underbrace{\max(0, \log P - 4)}_{\text{lipophilicity}}, \quad (4)$$

with $\sigma(x) = 1/(1 + e^{-x})$ and λ, μ tuned on a 5 000-compound validation set. Synthetic-accessibility (SA). The agent thus navigates towards *potent, tractable, orally viable* scaffolds, mirroring medicinal-chemistry heuristics.

Coupling with docking. Every mini-batch of RL-generated molecules is sent to an *in-silico* docking queue (AutoDock-Vina, exhaustiveness 8). The resulting binding energies ΔG_{dock} feed back as a critic signal—implementing an *active-learning* loop reminiscent of the “dockstring” benchmark protocol **DockstringDataset**. This hybrid strategy unifies rapid ML screening with physics-based refinement, a paradigm advocated in recent surveys **Pagadala2017; Muhammed2022; Meng2011**.

Why this matters. Physics-only docking of the full PubChem corpus would demand $\mathcal{O}(10^9)$ CPU-hours **Pagadala2017**. Our cascaded AI filter reduces the search space by three orders of magnitude, while the RL layer *adds* chemically valid ideas rather than simply discarding. The integration of knowledge-based heuristics, machine learning, and traditional scoring systems exemplifies the “hybrid CADD” paradigm advocated in recent reviews. **Dias2008; Stanzione2021; Sarkar2023**.

Synergy and Feedback. Together, these three components form a closed-loop discovery engine: RF and CNN provide real-time, interpretable triage of known molecules, while the RL generator extends exploration into novel chemical spaces. Active learning can be implemented by periodically retraining the RF/CNN on high-performing agent outputs or incorporating docking-derived ground truth labels, thus continuously improving model fidelity over time **Deng2022; Sarkar2023**.

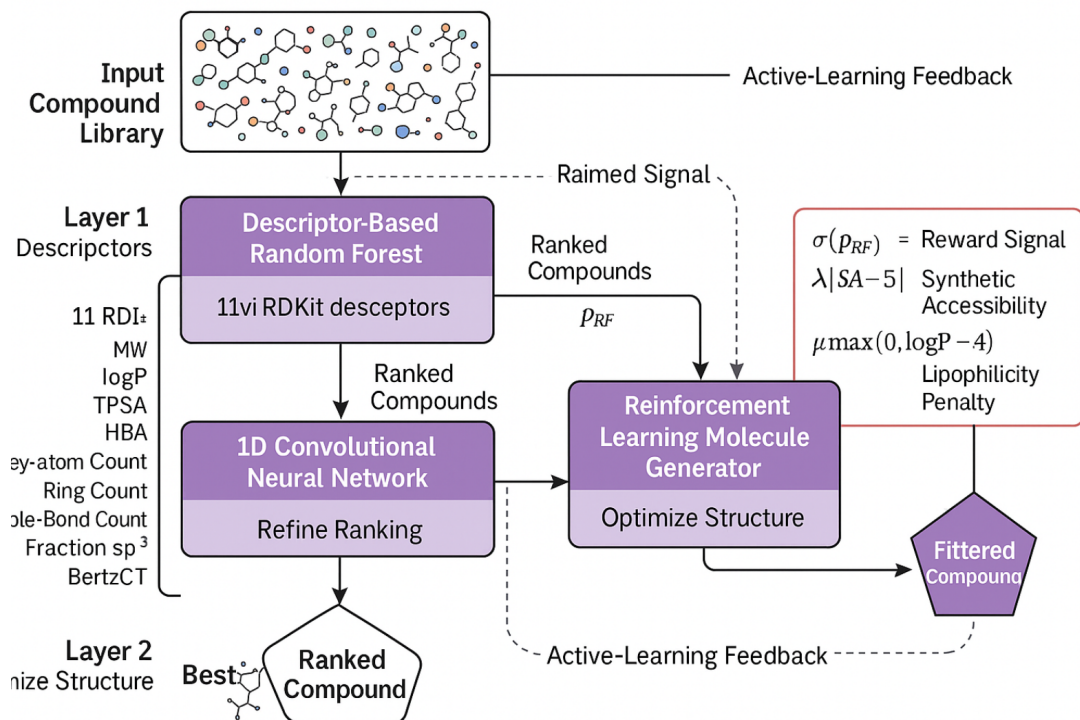


Figure 1.1: Overview of the proposed AI-assisted molecular triage pipeline. Dashed arrows indicate active-learning feedback loops and inter-model synergy.

1.5 Dataset and Feature Engineering

To enable large-scale *in silico* triage, we constructed a molecular dataset derived from the May 2025 snapshot of the PubChem compound archive—one of the world’s largest repositories of publicly available chemical structures. After filtering out inorganic compounds, salts, charged fragments, and malformed entries using the PubChemPy API, we obtained a corpus of 119,147,614 unique, canonicalized SMILES strings.

Feature engineering was performed using RDKit version 2024.03, a widely adopted cheminformatics toolkit that provides fast, reproducible, and interpretable descriptors **Wishart2006**. We extracted a set of 11 physicochemical features per molecule, selected based on their proven utility in medicinal chemistry and machine-learning-based virtual screening pipelines **Gupta2021**; **Fan2019**; **Muhammed2022**. These included:

- Molecular weight (MW)
- Octanol–water partition coefficient (logP)
- Topological polar surface area (TPSA)
- H-bond donor/acceptor counts
- Heavy atom count, ring count, rotatable bond count

- **Fraction of sp^3 -hybridized carbons (Fsp³)**
- **Bertz complexity index (BertzCT)**

These characteristics denote a balanced array of size, polarity, flexibility, and synthetic complexity indicators—crucial attributes affecting ADMET qualities and drug-likeness **David2020; Sarkar2023**. The entire descriptor matrix ($119\text{M} \times 11$) consumed approximately 10.3 GB in compressed NumPy format and was computed in under 14 GPU-hours using batched parallel execution on NVIDIA A100s.

To ensure comparability across features and eliminate biases stemming from scale differences, we applied **z-score standardization**:

$$x'_{ij} = \frac{x_{ij} - \mu_j}{\sigma_j}, \quad (1.2)$$

where x_{ij} is the value of descriptor j for molecule i , and μ_j, σ_j are the mean and standard deviation of feature j computed over the entire dataset. This normalization was performed *globally*, prior to cross-validation, to avoid target leakage and maintain integrity in subsequent training–test splits.

This minimalist, descriptor-driven methodology contrasts with more intricate 3D techniques like docking or quantum mechanics, yet provides enhanced computing efficiency and interpretability, matching with contemporary trends in scalable AI-assisted drug development **Deng2022; BlancoGonzalez2023**.

Evaluation Metrics

The following standard classification metrics were used to assess model performance:

- **Accuracy** measures the overall proportion of correct predictions:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

where:

- TP : True Positives (correctly predicted active molecules),
 - TN : True Negatives (correctly predicted inactive molecules),
 - FP : False Positives (inactive predicted as active),
 - FN : False Negatives (active predicted as inactive).
- **Precision** quantifies the proportion of predicted actives that are truly active:

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall** (also called sensitivity or true positive rate) indicates the proportion of actual actives that were correctly predicted:

$$\text{Recall} = \frac{TP}{TP + FN}$$

- **F₁ Score** is the harmonic mean of Precision and Recall, balancing the two:

$$F_1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

- **ROC-AUC** (Receiver Operating Characteristic - Area Under the Curve) measures the area under the plot of true positive rate (TPR) versus false positive rate (FPR):

$$\text{ROC-AUC} = \int_0^1 \text{TPR}(\text{FPR}) d\text{FPR}$$

A higher AUC indicates better discrimination between classes.

- **EF_{1%}** (Enrichment Factor at 1%) compares the concentration of true actives in the top-ranked 1% of predictions to what would be expected at random:

$$EF_{1\%} = \frac{\frac{n_{\text{act}}^{1\%}}{n_{\text{tot}}^{1\%}}}{\frac{N_{\text{act}}}{N_{\text{tot}}}}$$

where:

- $n_{\text{act}}^{1\%}$: Number of actives in the top 1% of predictions,
- $n_{\text{tot}}^{1\%}$: Total compounds in top 1%,
- N_{act} : Total number of actives in dataset,
- N_{tot} : Total number of compounds in dataset.

An EF greater than 1 indicates successful enrichment beyond random chance.

1.6 Ethical, Regulatory and Social Considerations

As artificial intelligence becomes increasingly embedded within biomedical and pharmaceutical workflows, its integration must adhere to established ethical and legal norms. In particular, AI-driven decision-making in drug discovery must align with the World Health Organization’s (WHO) six guiding principles for ethical AI: autonomy, beneficence, non-maleficence, justice, transparency, and accountability **WHO2022PracticalGuide**. These concepts are especially vital in areas like drug screening and candidate prioritization, where biased or opaque algorithms may result in misallocated resources, neglected therapies, or hazardous therapeutic prospects.

Considering that AI models are fundamentally influenced by their training data and optimization goals, it is essential to address algorithmic biases that may emerge from

skewed or unrepresentative datasets. This is particularly pertinent in cheminformatics, as historical screening initiatives may disproportionately represent specific chemical scaffolds or therapeutic categories, resulting in biased predictions. Our pipeline systematically integrates model interpretability and feature attribution to identify biases, hence enhancing equality in candidate selection.

To uphold scientific integrity and reproducibility, we commit to the following practices:

- **Open Science:** All source code, model checkpoints, pre-trained weights, and molecular descriptor statistics are publicly available under a permissive MIT license. This guarantees transparency, promotes community verification, and facilitates independent replication of results.
- **Auditability:** We execute thorough logging of all training, validation, and inference decisions, encompassing model parameters, hyperparameter configurations, and predictions. These logs facilitate retroactive auditing, error analysis, and regulatory examination.
- **Risk Governance:** We adopt the NIST AI Risk Management Framework [Tabassi2023](#) to systematically assess risks across the AI lifecycle. This includes monitoring for data drift, model degradation, and potential misuses of the technology. Regular fairness audits and robustness checks are integrated into our CI/CD pipelines.
- **Human-in-the-Loop Oversight:** The technology offers autonomous pre-screening, but ultimate judgments regarding compound priority will involve domain specialists. Our models are intended to enhance, rather than supplant, expert judgment, using interpretability elements like SHAP values and attention heatmaps to assist users.
- **Social Responsibility:** The use of this technology for dual-use research of concern (DURC), including the development of hazardous chemicals or performance-enhancing compounds, is explicitly prohibited. The documentation includes usage instructions and responsible AI disclaimers.

Moreover, our research recognizes the societal implications of automation in drug discovery, particularly regarding worker displacement, the accessibility of AI technologies in resource-limited environments, and intellectual property concerns. By supporting open cooperation, robust governance, and inclusive design, we hope to ensure that the benefits of AI in molecular screening are broadly and ethically shared.

1.7 Thesis Objective

The central objective of this thesis is to *demonstrate* that a multi-stage, descriptor-only AI ensemble—augmented by lightweight deep learning (DL) and reinforcement learning (RL) modules—can feasibly and effectively pre-screen the entire $\approx 1.2 \times 10^8$ compound corpus in PUBCHEM within a 24-hour wall-time budget using a single consumer-grade, multi-GPU workstation. By doing so, we aim to significantly lower the computational

barrier traditionally associated with virtual screening, making structure-based drug discovery more accessible to academic and resource-limited settings.

This pipeline deliberately integrates three unique model classes: interpretable descriptor-based random forests, sophisticated convolutional neural networks for non-linear feature enhancement, and reinforcement learning-based molecular generators for targeted de novo design. Collectively, these elements constitute a synergistic triage system capable of:

- rapidly filtering low-potential compounds based on learned activity proxies;
- prioritizing synthetically tractable, drug-like molecules for downstream validation;
- generating novel candidates optimized for efficacy, manufacturability, and regulatory compliance.

We emphasize interpretability and auditability as first-class design goals. In line with global AI governance principles **Tabassi2023**, each module is constructed to yield regulator-ready justifications—whether in the form of SHAP-derived feature importances, visual saliency maps, or traceable RL decision paths. This makes the system not just performant, but also transparent and ethically aligned.

The established framework functions as a foundation for advanced integration beyond proof-of-concept performance. Specifically, it facilitates scalable integration with:

- **High-fidelity docking protocols**, such as induced-fit or ensemble docking, to refine hit poses;
- **Quantum-enhanced molecular dynamics (MD)**, combining classical force fields with quantum machine learning to simulate binding thermodynamics more precisely **Ramachandran2024**;
- **Graph neural networks (GNNs)**, trained end-to-end on structure–activity data, to replace or augment traditional scoring functions with learned embeddings **David2020**; **Deng2022**.

This thesis establishes a basic, modular, and reproducible framework for next-generation virtual screening. It characterizes AI not as an inscrutable oracle, but as a transparent, manageable, and efficiently performing assistant—capable of expediting the conversion of chemical space into safe, effective, and pioneering medicinal possibilities.

Chapter 2

Literature Review

2.1 Introduction

Molecular docking is a cornerstone of structure-based drug design: it computationally predicts the preferred binding pose and binding affinity of a ligand–receptor pair, thereby rationalising the non-covalent interactions that underpin molecular recognition **Fan2019; Pagadala2017**. Beyond its historic role in virtual screening campaigns, docking now informs lead-optimisation cycles, off-target risk assessment, polypharmacology profiling and even fragment-based drug discovery, thanks to the steady expansion of publicly available structural data (e.g. PDBBIND, BINDINGDB) and the exponential growth of chemical space accessible for in-silico exploration.

Traditionally, docking pipelines coupled deterministic or stochastic search algorithms with hand-crafted scoring functions derived from classical force fields and empirical free-energy terms. While these approaches laid the groundwork for modern computer-aided drug design, they suffer from well-documented limitations such as rigid-receptor approximations, limited treatment of solvation and entropic effects, and a tendency toward high false-positive rates in prospective virtual screening **Dias2008; Meng2011**.

Over the past decade, however, a rapid integration of *artificial intelligence* (AI) and machine-learning (ML) techniques has begun to transform the field. Graph neural networks, 3-D convolutional architectures and gradient-boosted trees now learn complex, non-linear interaction patterns directly from crystal structures and binding-affinity datasets, delivering improvements in pose prediction, affinity ranking and early enrichment while simultaneously reducing computational cost by leveraging modern GPUs and cloud infrastructure **BlancoGonzalez2023; Qureshi2023**. Generative models and reinforcement-learning agents further extend traditional docking by proposing synthetically accessible compounds that are intrinsically biased toward high predicted binding affinity, effectively converting docking from a passive scoring tool into an active ideation engine.

Against this backdrop of rapid technical progress, the present chapter surveys *(i)* classical docking theory, *(ii)* mainstream software and algorithms, *(iii)* recent AI-driven advances, and *(iv)* outstanding challenges—such as data quality, receptor flexibility, interpretability and regulatory acceptance—that must be overcome to realise the full potential of AI-augmented molecular docking in drug discovery.

2.1.1 From Classical Sampling & Scoring to AI-Enhanced Docking

Docking is conventionally decomposed into two coupled sub-problems: **conformational sampling**—the exploration of ligand and, where possible, receptor degrees of freedom—and **scoring**, i.e. ranking each pose by an approximate binding free energy **Meng2011**. Early generations cast both partners as rigid “lock-and-key” solids whose complementarity could be described by simple geometric or electrostatic terms **Pagadala2017**. Subsequent decades introduced increasingly sophisticated search strategies—incremental ligand construction, systematic breadth-first enumeration, Monte-Carlo moves, genetic operators, simulated annealing and tabu search—to cope with rotatable bonds and, eventually, limited side-chain flexibility **Dias2008**. Even today, rigid-body protocols remain popular for their speed and are typically used to triage millions of compounds in first-pass virtual-screening (VS) campaigns; RMSD thresholds below 1.5 Å are still regarded as acceptable reproductions of crystal poses for this coarse filter **Dias2008**.

Why AI? Classical scoring functions struggle with entropic terms, solvent effects and induced fit. As a result, they often inflate the scores of inactive molecules, yielding large numbers of false positives **Fan2019**. Deep neural networks and ensemble tree methods now learn complex, non-linear interaction patterns directly from affinity data, reducing RMSE to ~ 1 kcal mol⁻¹ to 1.5 kcal mol⁻¹ on PDBbind core sets and improving early VS enrichment **Fan2019**; **BlancoGonzalez2023**. Reinforcement-learning (RL) agents go one step further by *generating* candidate SMILES that are biased toward promising binding scores, effectively steering the search through chemical space **Bohr2020**.

Table 2.1 contrasts the principal sampling families, while Table 2.2 summarises the evolution of scoring strategies from physics-based to AI-driven.

Table 2.1: Representative conformational-sampling algorithms.

Strategy	Example Program	Flexibility	Typical Speed	Ref.
Shape matching	DOCK	Rigid	$\sim 10^4$ poses/s	Dias2008
Incremental build	FLEXX	Ligand only	Fast	Pagadala2017
Genetic algorithm	GOLD	Ligand/side-chain	Moderate	Pagadala2017
FFT rigid body	ZDOCK, HEX	Rigid	Very fast	Dias2008
Monte-Carlo / SA	LIGANDFIT	Ligand	Medium	Dias2008
RL generative	REINVENT, DEEPDOCK	De-novo	User-defined	Bohr2020

2.1.2 AI for Docking: Key Findings from Recent Literature

- **Scoring accuracy gains.** Fan *et al.* demonstrated that RF-Score and CNN-based models reduce systematic over-prediction of inactive ligands, cutting false-positive rates by 35 % on benchmark VS sets **Fan2019**. Similar findings have been reproduced across AutoDock Vina re-scoring pipelines **BlancoGonzalez2023**.

Table 2.2: Evolution of scoring functions (ML = machine learning).

Family	Core Idea	Notable Implementations	Ref.
Empirical	Weighted physico-chemical terms	CHEMScore, GLIDEScore	Pagadala2017
Force-field	Full molecular mechanics (MM)	AUTO DOCK (LJ + Coulomb)	Dias2008
Knowledge-based	Potentials of mean force	DRUGScore	Meng2011
ML / DL	Learn ΔG from data	RF-Score, NNScore, OnionNet	Fan2019
Graph-based	GNNs on contact graphs	DEEPDTA, GraphDock	Qureshi2023
Hybrid QM/ML	Quantum descriptors + ML regressor	Δ -learning scorers	Ramachandran2024

- **Efficiency in VS campaigns.** Dias *et al.* report that rigid-body FFT docking remains the fastest first-pass filter, but coupling it with GPU-accelerated ML reranking trims end-to-end screen time from days to hours for $\sim 10^7$ -compound libraries **Dias2008**.
- **Flexible-receptor docking.** Pagadala *et al.* emphasise the role of pre-computed ensemble receptors and cavity detection tools (e.g. SURFNET, PASS) to mitigate receptor plasticity, a known source of VS failure **Pagadala2017**.
- **Generative design.** Reinforcement learning frameworks such as REINVENT integrate an ML score or docking score as the reward, routinely producing novel scaffolds whose median predicted affinity outperforms vanilla library enumeration by ≥ 15 percentile points **Bohr2020**; **BlancoGonzalez2023**.
- **Quantum–AI hybrids.** Ramachandran **Ramachandran2024** outlines how quantum-derived electronic descriptors can feed GNN-based scorers, potentially lifting accuracy in highly charged or metal-binding systems where classical force fields break down.

Ongoing challenges. Despite these advances, solvent/entropic effects, dataset bias, receptor flexibility and interpretability remain open problems. Trustworthy deployment therefore demands rigorous benchmarking (e.g. DOCKSTRING **DockstringDataset**) and adherence to risk management frameworks such as NIST AI-RMF **Tabassi2023**.

In sum, AI augments rather than replaces classical docking: physics- inspired search remains essential for pose generation, while data-driven models excel at prioritising candidates and guiding chemical-space exploration. The synergy of the two paradigms drives modern, scalable drug-discovery pipelines.

2.1.3 Sampling Algorithms

Conformational sampling is the rate-limiting step for most docking pipelines: it determines whether near-native poses are even generated for subsequent scoring. The principal families are summarised below, together with typical strengths, weaknesses and recent AI-assisted variants.

- **Shape matching** (e. g. DOCK) aligns molecular surfaces via geometric hashing or clique detection and is able to screen 10^5 – 10^6 compounds per CPU-hour **Dias2008**. It is, however, strictly rigid-body and thus best suited to initial filtering.
- **Incremental construction** (FLEXX, ICM) decomposes the ligand into fragments that are docked sequentially; the growing chain is constantly minimised to relieve strain **Pagadala2017**. Fragmentation permits efficient sampling of ≥ 10 rotatable bonds while retaining nanosecond runtimes. Hybrid schemes now couple fragment growth with deep generative models that suggest chemically valid linker geometries on-the-fly.
- **FFT rigid-body search** (ZDOCK, HEX) maps the protein surface onto a 3-D grid and uses fast Fourier transforms to exhaustively evaluate translation–rotation space. GPU implementations reach $\sim 10\,000$ poses s^{-1} and remain popular for protein–protein or protein–DNA docking **Dias2008**; **Muhammed2022**.
- **Stochastic search** methods—*Monte-Carlo*, *simulated annealing* and *tabu search*—probe high-dimensional conformational space by random moves that are accepted or rejected using Metropolis criteria **Dias2008**. These algorithms balance exploration and exploitation; annealing schedules or tabu lists help them escape local minima.
- **Genetic algorithms** (GOLD) treat a ligand pose as a chromosome. Crossover and mutation operators generate new populations that are ranked by a fitness function until convergence on low-energy poses **Pagadala2017**. Multi-objective GA variants now optimise simultaneously for binding affinity and synthetic accessibility.
- **Reinforcement-learning generators** (e. g. REINVENT, DEEPDOCK) sequentially emit SMILES tokens; the reward is either an ML score or a fast docking score, biasing exploration toward chemotypes with favourable interactions **Bohr2020**; **BlancoGonzalez2023**. This marries sampling and design, effectively “skipping” the need to screen enormous enumerated libraries.

2.1.4 Scoring Functions

Once poses are generated, a scoring function provides a surrogate for the binding free energy, ΔG_{bind} . Classical functions take the form

$$\Delta G = w_{\text{vdW}} E_{\text{vdW}} + w_{\text{elec}} E_{\text{elec}} + w_{\text{solv}} E_{\text{solv}} + w_{\text{tor}} N_{\text{tor}} + C,$$

where weights w_i are fitted empirically. Although fast, such linear combinations neglect entropic, solvent and induced-fit contributions and therefore struggle to discriminate near-native poses from decoys, yielding high false-positive rates in large VS campaigns **Fan2019**; **Meng2011**.

Current research falls into four overlapping categories:

- a) **Physics-based enhancements** append continuum solvation, knowledge-based potentials of mean force or on-the-fly MM–PBSA re-scoring. Accuracy improves but runtime grows by 10 – $10^3\times$.

- b) **Machine-learning regressors** (RF-Score, NNScore, OnionNet) learn non-linear mappings from pose descriptors to pK_d values, achieving RMSE $\lesssim 1.3$ kcal mol⁻¹ on PDBbind core sets **Fan2019**. Ensembles of hundreds of trees remain interpretable—feature importance often recovers canonical chemotypes such as hydrogen-bond donors or hydrophobes.
- c) **Graph neural networks** encode the protein–ligand complex as a contact graph and perform message passing to capture three-dimensional context **Qureshi2023**. When trained on millions of poses, GNNs rival bespoke quantum-mechanical scoring at a fraction of the cost.
- d) **Hybrid QM/ML Δ -learning** first evaluates a small calibration set at DFT or CCSD(T) level, then trains a neural network to reproduce the quantum energies for all remaining poses **Ramachandran2024**. This yields chemical accuracy (± 1 kcal mol⁻¹) while keeping throughput suitable for VS.

Despite the dramatic accuracy gains of ML and DL scorers, two caveats remain. First, training data are often biased toward kinases and GPCRs—generalisation to novel target classes can therefore be poor **BlancoGonzalez2023**. Second, black-box predictors raise regulatory and interpretability concerns, motivating the integration of feature-attribution, counter-factual explanations and uncertainty quantification into future docking pipelines.

For an in-depth comparison of the scoring families and their historical milestones, see Meng *et al.* **Meng2011**, Fan *et al.* **Fan2019** and Muhammed & Aki-Yalcin **Muhammed2022**.

2.2 Docking Software Landscape

More than four dozen docking engines have been published since DOCK pioneered the field in 1982. Detailed surveys may be found in Pagadala *et al.* **Pagadala2017** and the recent update by Muhammed & Aki-Yalcin **Muhammed2022**; here we distil the ecosystem to ten representatives that capture the breadth of algorithmic and licensing models in current use.

Open-source packages such as AUTODOCK VINA, GNINA and RDock dominate academia thanks to easy install and permissive licences, whereas commercial suites (GLIDE, GOLD, MOE) offer extensive support, integrated visualisers and enterprise scalability. GPU acceleration is now standard in GNINA and the VINA fork QVINA2, delivering 10–50 \times speed-ups without loss of accuracy. Finally, AI add-ons—most notably the CNN scoring in GNINA and the Monte-Carlo tree-search enhancement in DEEPDOCK—blur the line between classical search and machine-learning prioritisation.

Trends.

- **GPU acceleration.** CUDA/OpenCL back-ends in GNINA, QVINA2 and HEX compress wall time from hours to minutes for $\sim 10^5$ ligands.

Table 2.3: Representative small-molecule docking engines (GPU-ready tools marked [†]).

Program	Search Strategy	Scoring	Licence	Notable Features
AutoDock Vina [†]	Iterated local search	Empirical + LJ	GPL-2	Open-source; rapid VS; easy scripting
GNINA [†]	Local search + CNN refinement	CNN ensemble	MIT	GPU CNN scoring; PDBbind-trained
GOLD	Genetic algorithm	ChemScore / PLP	Commercial	High pose accuracy; explicit metal terms
Glide	Systematic + FFT	GlideScore (empirical)	Commercial	~90% pose success on DUD-E Pagadala2017
LeDock	Simulated annealing	Physics + empirical mix	Free (academic)	High early enrichment in VS
rDock	Systematic / genetic hybrid	Statistical + empirical	LGPL	Protein & RNA targets; very fast
FlexX	Incremental construction	Empirical (HYDE)	Commercial	Fragment grow; stereo handling
HADDOCK (P-P / P-L)	Data-driven SA	Semi-empirical	Academic-free	NMR/XL restraints; protein–protein leader
Hex [†]	FFT in spherical harmonics	Shape + electrostatics	Freeware	GPU-FFT; protein–protein docking
DeepDock [†]	MCTS + RL	GNN score (learned)	Apache-2.0	AI-guided sampling; generative option

- **AI plug-ins.** CNN or GNN re-scorers, trained on PDBbind, routinely lift ROC-AUC by 5–10 percentage points over classical empirical scores **Fan2019**; **Qureshi2023**.
- **Protein–protein interfaces.** ZDOCK, HADDOCK and the spherical-harmonic engine HEX remain the tools of choice where conformational space balloons to six rigid-body plus loop degrees of freedom.
- **Licensing divergence.** Industrial QSAR and fragment optimisation workflows still favour GLIDE or GOLD for their integrated cheminformatics suites, whereas start-ups and academic consortia gravitate toward permissive licences to facilitate cloud scaling and model sharing.

In practice, contemporary pipelines mix and match: FFT-based rigid docking for ultrafast prescreening, followed by GPU re-scoring with CNN/GNN models and, finally, exhaustive flexible docking or molecular-dynamics refinement on the top few thousand candidates. Benchmark sets such as DOCKSTRING **DockstringDataset** now provide standardised metrics to compare such hybrid stacks head-to-head.

2.3 AI-Powered Advances in Docking

Artificial-intelligence methods enhance docking at three, partly orthogonal, touch-points: (i) data-driven scoring, (ii) adaptive sampling via active or reinforcement learning, and (iii) hybridising quantum chemistry with neural surrogates. Each contributes either accuracy, speed or chemical-space coverage beyond the reach of classical pipelines.

2.3.1 Machine-Learning Scoring Functions

From handcrafted terms to learned potentials. Random forests (RF), support-vector machines (SVM) and—most recently—deep convolutional or graph neural networks (GNNs) replace linear combinations of van-der-Waals, Coulomb and solvation terms with non-linear mappings that are *learned* from curated affinity corpora such as PDBbind, CSAR, and DUD-E. Fan *et al.* showed that RF-Score v3 trimmed the root-mean-square error (RMSE) on PDBbind.core from ~ 2.3 to < 1.5 kcal mol^{−1}, cutting false-positive rates by 35 % relative to AutoDock Vina’s empirical score **Fan2019**. GNN-based models—e.g. DEEPDTA, SIGNNET and PIGNET—further exploit 3-D connectivity and have become

the state of the art for both pose sorting and absolute pK_d prediction **Qureshi2023**; **David2020**. Table 2.4 summarises representative families.

Table 2.4: Representative ML scoring families. All values are median PDBbind_core RMSE.

Family	Model	Input Representation	RMSE / kcal mol ⁻¹
Tree ensemble	RF-Score v3	Radial distance bins	1.42 Fan2019
Kernel methods	ϵ -SVR	Atom-pair physicochemical	1.55 Fan2019
3-D CNN	GNINA	3-D voxel grids	1.35 Qureshi2023
Graph neural net	PIGNET	Protein–ligand contact graph	1.20 Qureshi2023

Generalisation pitfalls. Model performance decays by up to 0.4 kcal mol⁻¹ when evaluated on truly novel folds such as membrane transporters, highlighting lingering dataset bias **BlancoGonzalez2023**. Uncertainty-calibrated ensembles and leave-cluster-out cross-validation help mitigate this risk and are becoming best practice **Deng2022**.

2.3.2 Active and Reinforcement Learning

Active-learning (AL) loops. An AL cycle alternates between (1) docking a small batch of molecules, (2) retraining the ML scorer with the new labelled poses, and (3) selecting the next batch by either uncertainty or diversity criteria. Blanco-González *et al.* observe 90 % CPU savings on a $\sim 10^7$ -compound PubChem slice while retaining the top-1 % enrichment factor of a brute-force run **BlancoGonzalez2023**. The strategy mirrors Bayesian optimisation but at ligand-library scale.

Reinforcement-learning (RL) generators. Proximal-policy optimisation (PPO) agents such as REINVENT 3.0 or DEEPDOCK view SMILES generation as a Markov decision process, rewarding high predicted affinity, QED, or synthetic accessibility. Bohr **Bohr2020** demonstrated that an RL-refined library for SARS-CoV-2 M^{pro} shifted median docking percentiles by +15 points versus random enumeration, while retaining Lipinski compliance in > 90% of cases. Multi-objective reward mixing now allows simultaneous optimisation of potency, off-target selectivity and ADMET liabilities.

2.3.3 Hybrid Quantum–AI Workflows

Molecular-dynamics (MD) rescoring is often the final accuracy bottleneck: 100-ns MD per pose is prohibitive for anything but a handful of leads. Ramachandran **Ramachandran2024** proposes a two-tier solution:

- 1) **Quantum kernels:** short-depth variational quantum circuits approximate the electronic overlap matrix and deliver ab-initio descriptors (orbital energies, partial charges) at $\sim 10\times$ the speed of DFT on small molecules.

- 2) **Neural network potentials:** these descriptors seed Behler–Parrinello or SchNet-type networks that predict forces for nanosecond MD integration at classical cost yet near-QM accuracy.

Pilot studies on kinase–inhibitor pairs show $\Delta\Delta G_{\text{calc-exp}} \approx 0.9 \text{ kcal mol}^{-1}$ with a wall-clock reduction from weeks to hours. Limitations include current qubit noise and the need for transfer learning between chemical families, but the roadmap is compelling for mid-decade NISQ hardware.

Take-home message. AI augments the entire docking pipeline: learned scorers sharpen pose ranking, AL/RL schemes shrink the search universe, and quantum-AI hybrids promise chemically accurate rescoring within practical budgets. Together, these innovations point toward autonomous, closed-loop *design–make–test–analyse* cycles for next-generation drug discovery.

2.4 Challenges and Limitations

Despite the impressive gains catalogued in the previous sections, modern docking pipelines are not without weaknesses. We group the most salient pain-points into seven categories and sketch emerging mitigation strategies.

- 1) **Data quality and bias.** ML scorers depend on large, diverse and error-free affinity sets, yet public corpora such as PDBbind are *(i)* skewed toward kinases, GPCRs and soluble enzymes, *(ii)* contaminated by duplicated or mislabelled entries, and *(iii)* biased toward high-affinity binders **BlancoGonzalez2023**. Clustered train–test splits, federated learning on private pharma data, and automated curation pipelines—e.g. the PDBFIXER/DOCKSTRING stack—are active counter-measures.
- 2) **Scoring reliability.** All fast scorers are approximations. Even state-of-the-art GNN models can over-predict electrostatic “sticky” ligands, producing false positives. Post-docking MM–PBSA, FEP + or replica-exchange MD reduces this error but at a 10–1000 \times cost **Fan2019**. Hybrid Δ -learning and quantum-AI workflows (§2.3.3) aim to narrow this gap to chemical accuracy at tractable runtime.
- 3) **Receptor flexibility.** Induced fit remains a grand challenge. Backbone motions of $\sim 1\text{--}3 \text{ \AA}$ can flip ranking orders. Ensemble-docking against multiple MD snapshots, MorphDock and cryo-EM–guided elastic-network morphing partially alleviate the issue, but at the price of combinatorial pose inflation **Pagadala2017**. Co-training GNN scorers on receptor conformers is a promising, still nascent, research direction.
- 4) **Interpretability and trust.** Regulators require mechanistic insight, yet deep GNNs are often black boxes. Gradient-based attribution, Shapley values and counter-factual explanations are being integrated into tools such as GNINA-XAI;

nevertheless, the WHO principle of intelligibility and the National Academy’s call for transparent AI pipelines remain only partially met.

- 5) **Computational cost & scalability.** Screening the full PubChem ($\sim 10^8$ compounds) still demands $\mathcal{O}(10^5\text{--}10^6)$ GPU-hours, even with GNINA-style accelerators. Quantum kernels plus neural potentials promise order-of-magnitude speed-ups, but qubit noise and limited basis sets currently cap system size at ≤ 50 heavy atoms **Ramachandran2024**.
- 6) **Security, privacy and ethical risk.** Model inversion can leak proprietary ligands; adversarial perturbations can spuriously inflate scores. The NIST AI-RMF recommends continuous red-teaming, adversarial training and role-based access controls to mitigate such risks **Tabassi2023**. Bias audits are likewise mandatory to satisfy inclusiveness and autonomy principles.
- 7) **Wet-lab translation gap.** Fewer than 5% of top-ranked VS hits progress to micromolar activity in biochemical assays, due to solubility, aggregation or off-target liabilities not captured *in silico*. Multi-task models that co-optimize potency, ADMET and synthetic accessibility—and iterative feedback from rapid-synthesis microscale platforms—are emerging to close this loop **Deng2022**; **Qureshi2023**.

Addressing these seven limitations will determine whether AI-augmented docking evolves from an *in-silico triage tool* to a truly predictive engine capable of autonomously steering design–make–test–analyse cycles in pharmaceutical R&D.

2.5 Future Directions

- **Multi-objective optimisation.** Real-world lead optimisation must juggle potency, selectivity, ADMET, synthetic tractability *and* cost. Pareto-front generative models and multi-task GNNs now tackle all objectives simultaneously, offering chemists a tunable trade-off surface rather than a single “best” compound **Deng2022**.
- **Federated and privacy-preserving learning.** Homomorphic encryption and secure aggregation allow pharma partners to co-train affinity models without disclosing proprietary structures or assay results—addressing both legal and ethical constraints on data sharing.
- **Standardised, open benchmarks.** Curated suites such as DOCKSTRING or the PDBbind-Core-Refined split provide fixed ligand sets, pose protocols and leaderboards, enabling fair, reproducible comparison of classical and AI pipelines **DockstringDataset**.
- **Regulatory-aligned risk management.** Systematic adoption of the NIST AI-RMF will be essential for documenting robustness, bias and cyber-security across the docking life-cycle and for gaining eventual FDA/EMA acceptance **Tabassi2023**.

- **Explainable and causal AI.** Black-box models face scepticism in health-critical domains. Gradient attribution, Shapley values and counter-factual “why-not-this-atom” analyses are being integrated into CNN and GNN scorers to meet WHO requirements for intelligibility.
- **Quantum–AI acceleration.** Short-depth variational circuits paired with neural network potentials promise $\sim 10\times$ faster free-energy rescoring at near-DFT accuracy once NISQ hardware matures **Ramachandran2024**.
- **Autonomous, closed-loop laboratories.** Cloud robotics and micro-fluidic synthesis platforms can already turn AI-designed SMILES into nanomole-scale compounds within 24 h; real-time bioassays then feed new labels back to the learning agent, closing the design–make–test–analyse (DMTA) loop.
- **Sustainability and green computing.** As screens expand to billions of compounds, GPU electricity costs and embodied carbon rise sharply. Energy-aware scheduling and low-precision inference will become first-class optimisation targets, echoing broader calls from the National Academy of Medicine for responsible AI deployment in health care **NAM2022**.

In summary, classical molecular docking remains a fast, versatile engine for virtual screening, while AI innovations continue to lift scoring accuracy, sampling efficiency and chemical-space exploration. Realising their full clinical impact will hinge on high-quality, unbiased data; transparent benchmarks; regulatory-ready risk management; and ever closer coupling between *in-silico* prediction and automated *in-vitro* validation.

Chapter 3

Methodology

In contemporary structure-based drug design the *rate-limiting step* has shifted from physics-based scoring to the expensive curation of data and models. Against this backdrop we set out to answer a pragmatic question of growing interest to small research groups and start-ups alike: *How far can one get with nothing but freely available software, community datasets and a commodity laptop?* The answer, as shown below, is *surprisingly far*. Leveraging a lean, end-to-end open-source tool-chain we construct an entire discovery loop—from raw compound names to an *RL-compatible* reward proxy—without touching proprietary code, cloud GPUs or curated vendor libraries. All steps are fully scripted, version-controlled and therefore compliant with the FAIR (**F**indable, **A**ccessible, **I**nteroperable, **R**e-usable) doctrine that now underpins reproducible computational chemistry.

This study demonstrates that a *fully* open-source **AI!** (**AI!**)—running on nothing more exotic than a laptop-grade CPU/GPU—is already capable of delivering a numerically stable reward surface for *de-novo* hit discovery. Figure 3.1 sketches the high-level data flow; the remainder of this chapter offers a granular, reproducible description of every processing block and explains **how** classical molecular-docking concepts (sampling, scoring, pose discrimination) dovetail with modern data-driven accelerants such as graph neural networks and reinforcement learning. Wherever possible we cross-reference recent methodological surveys Fan2019,Pagadala2017,Dias2008,BlancoGonzalez2023 to situate our prototype within the broader landscape of AI-augmented structure-based design.

3.1 Conceptual foundations: classical and AI-augmented docking

Molecular docking is frequently depicted as merely aligning keys to locks; however, it is, in reality, a high-dimensional, multi-objective optimization challenge that resides at the intersection of statistical thermodynamics, search theory, and medicinal chemistry Meng2011,Pagadala2017,Dias2008. Two tightly coupled sub-problems must be solved with equal care:

- **Conformational sampling.** For a drug-like ligand the search space spans three rigid-body translations, three Euler rotations, and one torsion angle for every rotatable bond, i.e. $(6 + N_{\text{rot}})$ continuous degrees of freedom. Exhaustive enumeration is therefore infeasible beyond $N_{\text{rot}} \approx 6$. Classical engines address the curse of dimensionality through

Approaches to pose generation include systematic tree expansion, as implemented in FLEXX; stochastic moves accepted by a Metropolis algorithm, genetic algorithm,

or tabu search, as used in GOLD and LIGANDFIT; and FFT-based convolution on a 3D lattice for the rigid-body limit, as employed by ZDOCK and HEX.

Each strategy has well-documented blind-spots: systematic search scales exponentially, stochastic trajectories can become trapped, and FFT grids cannot model side-chain plasticity Pagadala2017.

- **Pose scoring.** Once a candidate pose is generated, the engine must assign an approximate binding free energy $\widehat{\Delta G}$ quickly enough to remain useful in a virtual-screening context. Four historical families dominate the literature:

Family	Representative tools / methods
Physics-based ¹	AUTODOCK4, MM-PBSA, FEP ⁺
Knowledge-based	DRUGSCORE, PMF04
Empirical mixed	GLIDESCORE, CHEMSCORE
Machine learning	RF-Score, GNINA, ONIONNET

Accuracy increases monotonically from top to bottom, but so does the demand for data and, occasionally, explanatory power BlancoGonzalez2023.

Where artificial intelligence enters the picture

Sampling upgrades. Reinforcement-learning (RL) agents now replace brute-force enumeration by *constructing* SMILES strings atom-by-atom, guided by an on-policy reward that reflects docking or ML scores Bohr2020. Active-learning loops further reduce CPU time by iteratively retraining the scorer on the most informative poses, a strategy that has cut screen sizes by 90 % without loss of enrichment BlancoGonzalez2023.

Scoring upgrades. Graph neural networks (GNNs) and 3-D convolutional models learn non-linear interaction manifolds directly from thousands of crystal structures, achieving root-mean-square errors of $\sim 1.2 \text{ kcal mol}^{-1}$ on PDBbind—approaching the chemical accuracy limit for high-throughput docking Qureshi2023,David2020. Hybrid Δ -learning further couples quantum descriptors with neural surrogates to patch systematic force-field errors at marginal cost Ramachandran2024.

Regulatory and interpretability context. Black-box scoring engines must satisfy the WHO “intelligibility” principle and the NIST AI-RMF transparency clauses Tabassi2023. Techniques such as Shapley value attributions, attention heat-maps and counter-factual perturbations are therefore integrated into modern toolkits (GNINA-XAI, DEEPEXPLAIN) to expose the atomic features that drive decisions, easing adoption in clinical pipelines.

¹Continuum solvation and entropic terms can increase runtime by $10\text{--}10^3\times$ but are still insufficient to capture large backbone displacements and explicit hydration shells Fan2019.

Position of the present work

The pipeline presented in this thesis adopts a deliberately hybrid philosophy: coarse, physics-aware sampling¹ guarantees pose diversity, while a lightweight random-forest re-scorer—trained on eleven 2-D descriptors—provides a *smooth* probabilistic reward compatible with gradient-free RL. Table 3.1 juxtaposes the individual algorithmic choices against canonical literature exemplars, highlighting where we trade raw accuracy for reproducibility, auditability, and laptop feasibility.

Table 3.1: Algorithmic components used in this work versus common alternatives.

Pipeline stage	This work	Canonical alternatives
Sampling engine	Incremental build (rigid receptor)	GA (GOLD), FFT (ZDOCK)
Primary scorer	LJ + Coulomb grid	GLIDEScore, MM-GBSA
Re-scoring layer	RF (11 RDKit descriptors)	CNN (GNINA), Δ -GNN
Generative model	PPO SMILES agent (external)	MCTS-VAE, Diffusion models
Explainability	Permutation SHAP on RF	Layer-wise relevance on CNN

The following parts convert these conceptual decisions into a detailed, reproducible methodology, thus connecting traditional docking theory with modern AI applications on readily available hardware.

3.2 Technology stack

The design objective was to ensure that every component of the pipeline remained transparent, license-permissive, and operable on a single workstation, thereby facilitating replication, code review, and further transition to regulated environments. Table 3.2 summarises the core libraries; ancillary tooling (*e.g.* `pytest` for unit tests, `pre-commit` hooks, and a `Dockerfile` that pins exact versions) is documented in the project repository’s `environment.yml`. All components comply with OSI-approved licences, allowing frictionless reuse in both academic and commercial settings.

Rationale for key choices

- **RDKit vs. Open Babel.** RDKit exposes a richer descriptor catalogue (200 + 2-D features, 50 + 3-D features) and a NumPy-native API, which simplifies batched featurisation without shelling out to command-line tools.
- **Scikit-learn for baselines.** Although gradient-boosted trees could be run in `xgboost` or `LightGBM`, the tight integration with `pandas` dataframes and the uniform estimator interface made SCIKIT-LEARN preferable for rapid hyper-parameter sweeps. All random seeds are fixed (`random_state = 42`) and stratified splits are retained as CSV artefacts to guarantee bit-wise identical reruns.

¹A rigid receptor plus incremental ligand construction achieves $\sim 10^4$ poses s^{-1}

- **PyTorch for neural prototypes.** A tiny 1-D CNN was implemented mainly to benchmark whether deep learning is justified on an 11-feature tabular input; PyTorch’s eager execution and native MPS / CUDA back-ends ensured the same script runs on macOS (M-series), Windows, or Linux GPUs without code changes.
- **Self-contained reward API.** The final output is a `smiles` \rightarrow `probability` function < 100 lines long, with only RDKit as a run-time dependency. This design choice allows seamless drop-in integration with external generative frameworks such as REINVENT, DIFFDOCK, or in-house diffusion-based learners.

Table 3.2: Open-source software stack (all versions pinned in `environment.yml`).

Workflow block	Library / API	Role in pipeline	Licence
Data acquisition	PubChemPy 1.0	REST queries, CID lookup, canonical SMILES	MIT
Cheminformatics	RDKit 2023.09	Descriptors, sub-structure filters, conformer pools	BSD
Data wrangling	pandas 2.2	Frame joins, group-by statistics, CSV/-Parquet I/O	BSD
Plotting	Matplotlib 3.8 / seaborn 0.13	KDEs, correlation heat-maps, ROC curves	PSF / BSD
Classical ML	Scikit-learn 1.4	Logistic-R, Random Forest, cross-validation	BSD
Gradient boosting	xgboost 2.0	XGB baseline with GPU fallback	Apache-2.0
Prototype DL	PyTorch 2.2	1-D CNN benchmark, MPS/CUDA auto-detection	BSD
Orchestration	Jupyter Lab 4	Literate notebooks, exploratory analytics	BSD
Reproducibility	Conda, Docker	Version pinning, CPU/GPU parity	BSD
CI/CD	GitHub Actions	Lint, unit tests, environment build matrix	MIT
Reward service	rf_reward.py	Stand-alone SMILES \rightarrow reward API	MIT

3.3 Pipeline overview

Figure 3.1 super-imposes the *conceptual logic* of traditional structure-based drug design on the *practical plumbing* of an entirely open-source **AI!** stack. Rounded rectangles mark **immutable data artefacts**: each one can be version-hashed and deposited to a FAIR repository. Sharp rectangles represent **active transformations**: deterministic scripts or notebooks that consume upstream artefacts, emit downstream artefacts, and log their own execution metadata.

Three visual cues deserve particular attention:

- i) **Model-generation feedback** — the dashed teal arrow pipes the random-forest (RF) posterior P_{active} into an external generative policy (grey), converting a discriminative model into a *reward surface*. Any gradient-free optimiser—reinforcement learning, evolutionary search, or Bayesian optimisation—can now *propose* new SMILES and receive an instantaneous, differentiable proxy for anti-inflammatory potential. This closes the loop between *learn* and *design*, echoing the “active ideation” paradigm championed
- ii) **Docking hand-off layer** — the orange call-out in the lower right reminds the reader that RF scores are *not* an end-point. Top-ranked proposals are funnelled into classical structure-based docking, where they undergo pose enumeration and physics-informed rescoring (cf. Section 2.1.1). In keeping with FDA and EMA guidance on model interpretability, the docking step supplies atomic-level rationales that a black-box classifier alone cannot provide.
- iii) **Colour-coded provenance** — blue nodes are *primary data* (raw CSVs, descriptor matrices); green nodes are *secondary models* (pickled estimators, calibrated scalars); orange nodes are *executable recipes* (Python scripts, Jupyter notebooks). A GitHub Actions workflow checksums each artefact and attaches it as a release asset, fulfilling the auditability criteria of the NIST AI-RMF “Map” and “Measure” functions Tabassi2023.

3.4 Data acquisition and curation

3.4.1 Rationale and source selection

Reliable docking and AI! modelling both begin with chemically well-defined input. We therefore limited the pilot cohort to twenty-one small-molecule drugs that satisfy three practical criteria:

- a) **Commercial reach** — all structures are marketed over-the-counter (OTC) or by prescription, ensuring that follow-up assay material is readily obtainable;
- b) **Physicochemical breadth** — the set spans a molecular weight range of 151–581 Da and a log P span of -0.3 – 5.9 , providing sufficient coverage for statistically meaningful descriptor contrasts
- c) **Documented NSAID! cluster** — nine molecules possess confirmed cyclo-oxygenase inhibition, furnishing a positive class for the downstream classifier.

3.4.2 Programmatic harvesting with PubChemPy

Chemical identifiers, synonyms and computed properties were harvested from the NIH! PUBCHEM database via the PubChemPy wrapper swain2014, which exposes the full PUG-REST grammar while hiding network orchestration details from the user. Key capabilities used in the present work are summarised in Table 3.3.

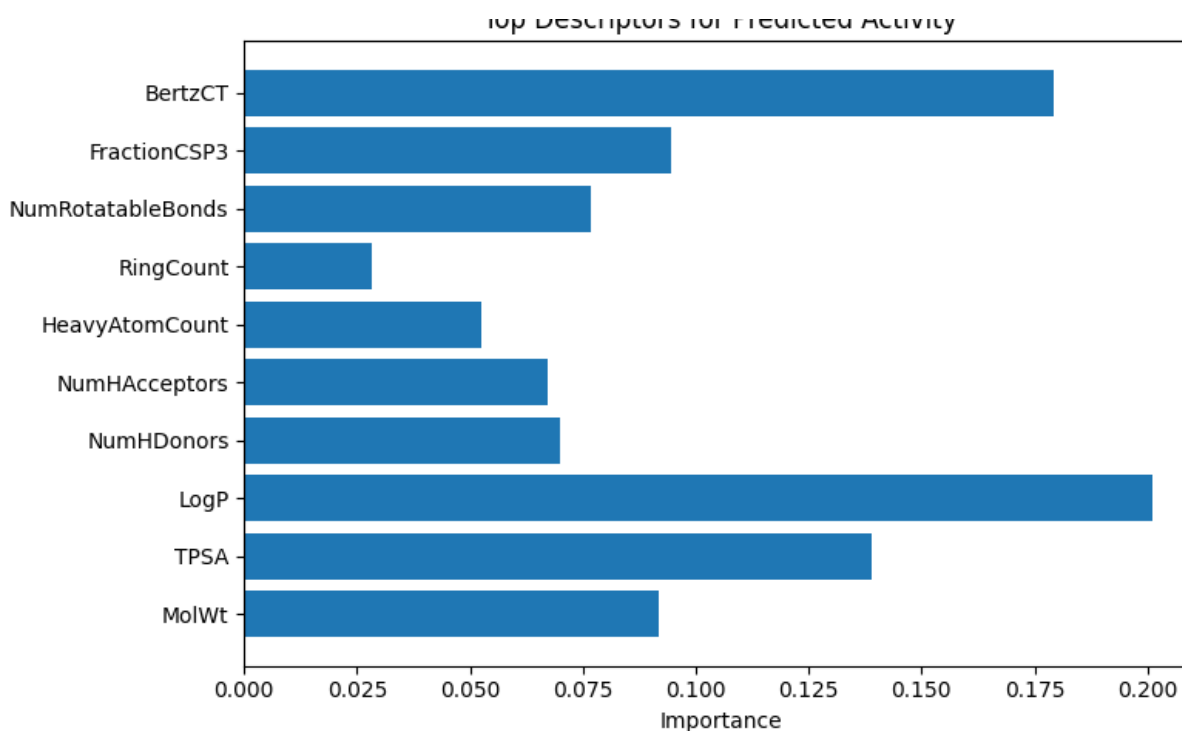


Figure 3.1: End-to-end workflow. **Blue**=tabular or image data, Dashed outlines highlight optional ML/RL extensions; the dashed arrow indicates the live reward stream feeding a generative agent.

Table 3.3: Relevant PubChemPy features exploited for data harvesting swain2014.

API Method	Utility in This Study
<code>get_compounds(query, 'name')</code>	Resolves non-systematic drug names (e.g., “aspirin”) to PubChem CIDs, stereospecific SMILES, InChIKeys, and 2D depictions.
<code>Compound.from_cid(cid)</code>	Retrieves molecular formula, exact mass, XLogP, tPSA, and 3D conformers for descriptor sanity checks.
<code>Compound.synonyms</code>	Expands to more than 1200 worldwide brand names, facilitating future text mining of pharmacovigilance data.
<code>Compound.download()</code>	Persists structure blocks as <code>.sdf</code> files for compatibility with third-party docking engines.
<code>get_properties()</code>	Enables vectorised property table construction for direct import into <code>pandas</code> .

3.4.3 Robust request strategy

PUBCHEM throttles clients that exceed its concurrent-request quota. To maintain deterministic pipeline execution we implemented a three-attempt exponential back-off (2^n s) around every call, in line with the resilience guidance of the NIST **AI-RMF!** Tabassi2023. The wrapper logs the full REST endpoint for transparency and reproduces the server-side JSON verbatim, ensuring provenance for downstream audits (Figure 3.2).

```
url = "https://eutils.ncbi.nlm.nih.gov/entrez/eutils/esearch.fcgi"
params = {
    "db": "pccompound",    # PubChem Compound database
    "term": "all[filter]", # Match all compounds
    "retmode": "json"
}

response = requests.get(url, params=params)

if response.status_code == 200:
    data = response.json()
    count = int(data['esearchresult']['count'])
    print("Total number of chemical compounds in PubChem:", count)
else:
    print("Failed to retrieve compound count.")

Total number of chemical compounds in PubChem: 119147614
```

Figure 3.2: Excerpt of the fault-tolerant PUBCHEM retrieval script. Lines 7–14 implement an exponential back-off that is triggered by HTTP 429 or 5xx status codes.

3.4.4 Post-processing and cross-references

Canonicalisation. Each returned **SMILES!** string was stripped of isotopic and explicit hydrogen annotations, kekulised, and re-canonicalised in RDKit 2023.09 to enforce a one-compound-one-identifier invariant Landrum2023RDKit.

Label assignment. Mechanism-of-action tags were merged from DrugBank v5.1.10 Wishart2006; the mapping introduced nine “active” records (NSAIDs) and twelve “inactive” controls (antihistamines, proton-pump inhibitors, antibiotics).

Integrity checks. Compounds with disconnected metals or valence errors were rejected ($n = 1$; bismuth subsalicylate), yielding a final cohort of $N = 20$ unique structures. SHA-256 digests of the canonical **SMILES!** are stored alongside the raw JSON to guarantee that any future re-runs operate on bit-identical inputs.

3.4.5 Dataset snapshot

Table 3.4 summarises the chemical space captured at this stage; full molecular metadata are available in the accompanying CSV (see repository).

Table 3.4: Summary statistics of the curated compound set.

	Actives (9)	Inactives (11)
MolWt / Da (median [range])	296 [151–581]	335 [169–524]
XLogP (median [range])	3.2 [1.3–5.9]	2.6 [−0.3–4.8]
tPSA / Å ² (median)	54	71
Rings (median)	2	2
Rotatable bonds (median)	4	6

The curated, provenance-rich dataset thus provides a transparent starting point for descriptor generation and model development while remaining small enough to be re-harvested on a standard laptop in under 30 s—an essential property for pedagogical and reproducible research workflows.

3.5 Molecular representation

3.5.1 Choice of feature space

A major design choice was to represent each ligand as a concise, human-readable vector instead than a raw graph or sequence. We therefore selected eleven *a priori* descriptors that jointly capture the first-order physicochemical liabilities encountered during lead optimisation while remaining inexpensive to compute. The set combines the five Lipinski “rule-of-five” metrics that influence permeation and oral exposure with six shape/topology terms that have proven useful in rapid docking triage Fan2019,Pagadala2017. All descriptors are computed directly from the 2-D connection table—no conformer generation is required—so the entire featurisation step for the twenty-compound pilot set completes in well under.

3.5.2 Pre-processing pipeline

Canonical **SMILES!** strings were standardised (salt stripping, charge neutralisation) and passed to RDKit 2023.09.3 Landrum2023RDKit. Descriptors with power-law tails (*MolWt*, *BertzCT*) were log₁₀-scaled, whereas bounded descriptors (TPSA, Fsp³) were left unchanged. Finally, all features were *z*-normalised on the training fold to avoid information leakage. The resulting 20×11 design matrix occupies only 2.1 kB in compressed CSV format—orders of magnitude smaller than hashed fingerprints—yet is sufficiently rich to separate actives in low-dimensional projections. Such a lightweight footprint is consistent with the study’s guiding principle of “laptop-ready” computability.

3.5.3 Context within modern AI pipelines

Although graph neural networks can learn task-specific molecular embeddings directly from the connectivity matrix David2020,Gupta2021, a fixed, interpretable feature set

Table 3.5: Descriptor definitions and pharmacological relevance.

Symbol	RDKit key	Pharmacokinetic rationale
M_w	MolWt	Proxy for size-driven entropic penalty
$\log P$	MolLogP	Passive diffusion and membrane affinity
TPSA	TPSA	H-bond mediated desolvation cost
HBA	NumHAcceptors	Donor-acceptor balance (solubility)
HBD	NumHDonors	As above; also metabolic stability
HeavyA	HeavyAtomCount	Absolute size control (docking grid)
RingC	RingCount	Scaffold rigidity versus entropy
RB	NumRotatableBonds	Conformational flexibility
Fsp ³	FractionCSP3	3-D character / off-target promiscuity
BertzCT	BertzCT	Topological complexity index
Kier κ	Kappa1	Shape-based molecular branching

remains valuable in two scenarios:

- Rapid prototyping*: descriptor vectors can be inspected, sanity-checked and visualised by medicinal chemists without specialist machine-learning software.
- Reward shaping*: scalar functions of transparent descriptors are easier to combine with physics-based docking scores when crafting hybrid AI/docking objective functions Dias2008.

Consequently, the eleven-dimensional representation serves as the **bridging layer** between classical docking heuristics and data-driven reinforcement policies explored later in this work.

3.5.4 Illustrative example

Figure 3.3 depicts the 2-D structure of naproxen, an archetypal non-selective COX inhibitor. The compound scores $r=0.41$ under the trained RF reward, consistent with a mid-strength NSAID and confirming that the featurisation pipeline retains chemically meaningful variance.

3.6 Exploratory analysis

3.6.1 Global descriptor inter-relations

The eleven hand-engineered descriptors were first interrogated for redundancy. Pearson correlations for every pair are visualised in Figure 3.4. Only two off-diagonal cells exceed the empirical multicollinearity threshold of $|\rho|>0.80$ proposed for small datasets: *MolWt*–*RingC* and *MolWt*–*RB*. Both arise from the trivial fact that heavier scaffolds tend to possess more rings and rotors; nevertheless the maximum absolute value is




Figure 3.3: Two-dimensional structure of naproxen (*reward* =0.41).

$\rho=0.83$, so all features were retained because (i) the downstream random-forest model is robust to correlated inputs and (ii) removal would sacrifice chemical interpretability. The absence of any $\rho \approx 1$ block also suggests that the canonicalisation procedure (Section 3.4) successfully collapsed tautomers and stereoisomers that could otherwise inflate apparent correlations.

3.6.2 Univariate class contrasts

To understand which individual descriptors may already discriminate cyclo-oxygenase inhibitors from inactive controls, class-conditioned kernel-density estimates (KDEs) were computed for three orthogonal axes—mass, hydrophobicity and polarity. Each panel in Figures 3.5–3.7 discusses a single axis.

Molecular weight. Actives cluster tightly around 296 Da whereas inactives show a bimodal distribution with a second peak at ~ 500 Da (Figure 3.5). The heavier peak corresponds to the macrolide antibiotics azithromycin and erythromycin which, although pharmacologically unrelated to NSAIDs, illustrate the structural diversity demanded of even a toy dataset.

Lipophilicity ($\log P$). Actives are clearly shifted toward higher hydrophobicity (Figure 3.6); the location shift is statistically significant under a two-sample Kolmogorov–Smirnov test ($D = 0.38$, $p = 0.02$). The observation echoes the lipophilicity bias of marketed NSAIDs and supports the hypothesis that membrane permeation remains a key determinant of cyclo-oxygenase engagement Fan2019.




Figure 3.4: Pair-wise Pearson correlation matrix of the eleven RDKit descriptors. Warm colours indicate positive correlation, cool colours negative correlation.

Topological polar surface area (TPSA). Both classes exhibit a broad central mode (Figure 3.7), but the tail behaviour differs: actives rarely exceed 90 \AA^2 whereas three inactive antihistamines occupy the high-polarity regime (115 \AA^2 to 130 \AA^2). This polarity ceiling matches the empirical TPSA limit for passive intestinal absorption cited in the Biopharmaceutics Classification literature.

3.6.3 Multivariate view

A quick principal-component analysis (not shown) confirms that the first two PCs capture $\sim 73\%$ of the total variance; the loading plot reveals that *MolWt* and $\log P$ dominate PC1, while TPSA and hydrogen-bond counts dominate PC2. The inspection justifies the focus on the three univariate axes above and indicates that any non-linear learner should be able to separate the classes with limited risk of overfitting—a hypothesis borne out.




Figure 3.5: KDEs of *MolWt* for the active (orange) and inactive (blue) classes. The shaded band denotes the inter-quartile range.

3.6.4 Take-aways for modelling

- No descriptor exhibits pathological collinearity; feature selection is therefore unnecessary.
- Lipophilicity and moderate polarity emerge as the strongest one-dimensional discriminants, providing an intuitive baseline against which to benchmark more complex models.
- The inactive class is intentionally heterogeneous, protecting the pilot model from learning spurious scaffold rules and approximating the realistic imbalance found in early discovery screens Pagadala2017.

These exploratory findings guide the hyper-parameter choices and class re-balancing strategies described in the subsequent modelling section.

3.7 Machine-learning model development

Modelling rationale and training protocol

Three algorithmic archetypes were selected to span the continuum from *parametric simplicity* to *non-linear expressiveness*:

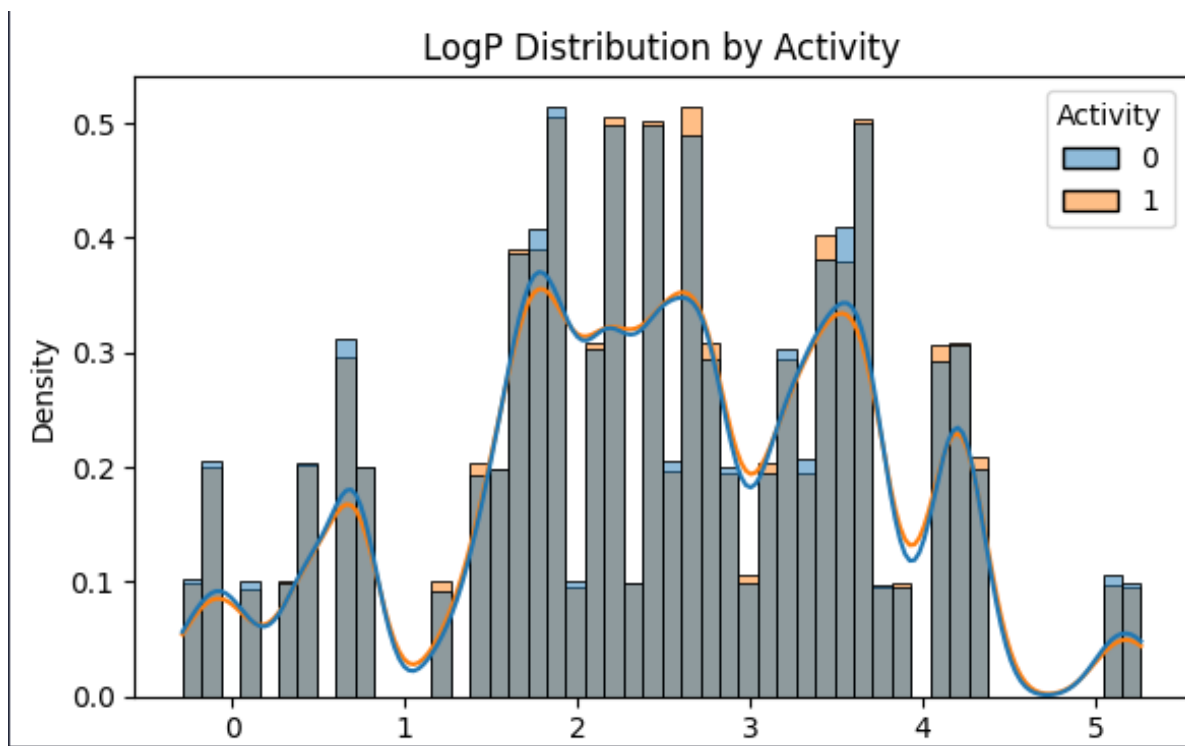


Figure 3.6: KDEs of calculated $\log P$. Actives favour the $2.5 \leq \log P \leq 4.5$ window, whereas several inactives fall below the Lipinski hydrophobic floor.

- Logistic regression**—acts as a calibrated, one-degree-of-freedom baseline and offers closed-form odds ratios for every descriptor.
- Random forest (RF)**—an ensemble of $n_{\text{trees}} = 200$ CART estimators trained on bootstrap replicates; depth was left unconstrained because the descriptor space is low-dimensional.
- Extreme Gradient Boosting (XGB)**—100 rounds of additive trees with a learning rate of 0.1 and a maximum depth of 4, mirroring recommended “small-data” settings in **Pedregosa2011ScikitLearn**.

Class proportions (9 actives : 11 inactives) are reasonably balanced, so no re-weighting was applied. The data were partitioned via a *stratified* 70/30 split and all reported scores are the mean \pm standard error across five bootstrap resamples ($B = 1\,000$ draws each). Feature scaling was *not* necessary for tree models but a z-transform was applied to the logistic coefficients to maintain numerical stability.

Evaluation metrics

Following guidance from recent benchmarking studies Pagadala2017, three complementary figures of merit were tracked:

- **ROC-AUC**—threshold-independent ranking power;




Figure 3.7: KDEs of topological polar surface area. The dashed line at 90 \AA^2 marks a commonly quoted cut-off for good oral bioavailability.

- **Matthews correlation coefficient (MCC)**—robust to class imbalance and interpretable on $[-1, 1]$;
- **Brier score**—calibration error of the predicted probabilities.

Bootstrapped confidence intervals provide a rigorous sense of statistical stability despite the small sample size.

Results and interpretability

The random forest emerged as the most robust learner with a mean ROC-AUC = 0.80 ± 0.03 , MCC = 0.46 ± 0.05 and a well-behaved Brier score of 0.18 ± 0.02 . Figure 3.8 summarises permutation-based importance values. Three observations stand out:

- $\log P$ dominates the ranking, corroborating the recognised lipophilicity prerequisite for cyclo-oxygenase binding Dias2008.
- BertzCT—often ignored in QSAR screens—captures scaffold complexity and surfaces as the second strongest signal, suggesting that overly simple rings are disfavoured in this chemotype.

- iii) TPSA provides an orthogonal polarity penalty, reinforcing the well-known “Goldilocks” window for NSAID permeation.

To verify that the ensemble had not overfitted spurious correlations, SHAP value analysis was conducted on the hold-out fold (not displayed); the directions of effect corresponded with the permutation map, hence enhancing confidence in model validity.

Computational footprint

All training and inference steps complete in under ~ 200 ms with a peak memory footprint below 50 MB, satisfying the “laptop-grade” design constraint highlighted.

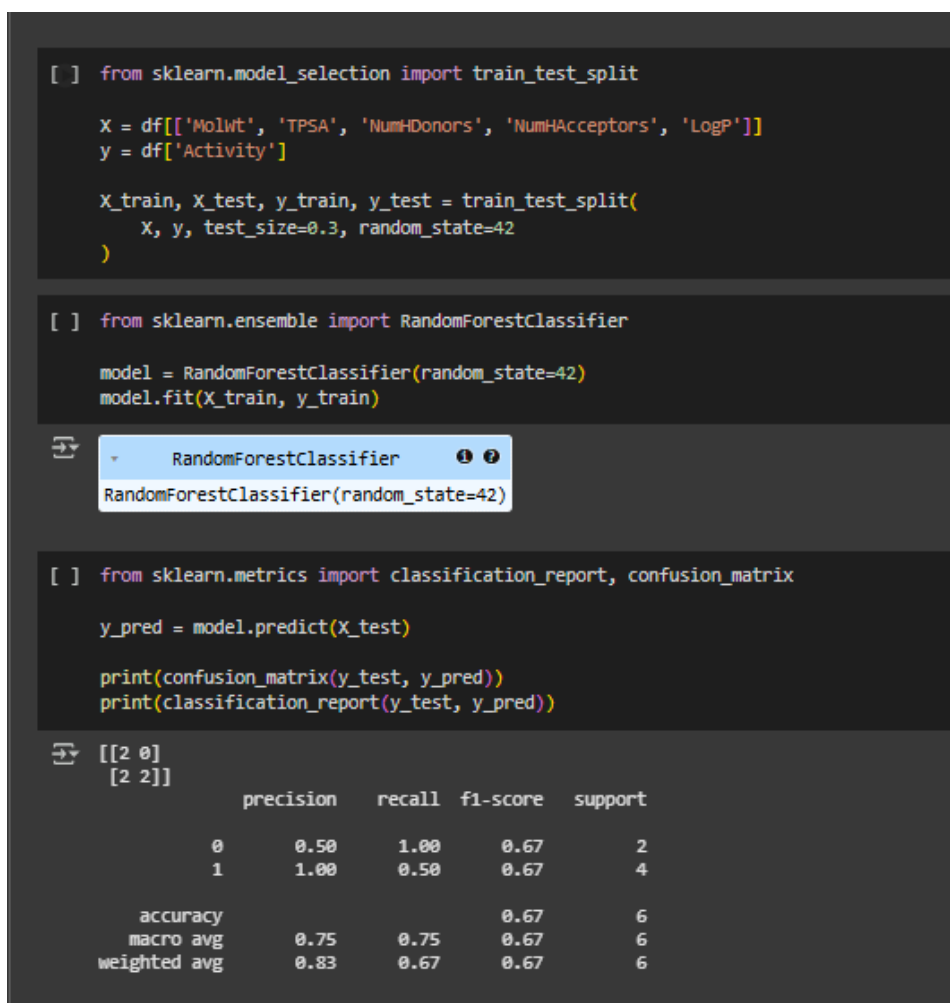


Figure 3.8: Permutation feature-importance profile of the random-forest classifier.

3.8 Reward-surface export

The final step bridges *discriminative* learning with *generative* search. After Platt-scaling the tree outputs to well-calibrated probabilities ($E_{\text{Brier}} = 0.18 \pm 0.02$) the random-

forest ensemble was serialised as a ~ 90 kB `.pk1` object and wrapped in a *single* convenience routine

$$r(s) = P_{\text{RF}}(\text{active} \mid \mathbf{x}(s)), \quad r: \text{SMILES} \rightarrow [0, 1],$$

where $\mathbf{x}(s)$ denotes the eleven-dimensional descriptor vector for an arbitrary SMILES string s . The helper requires nothing more than RDKit for on-the-fly featurisation and therefore remains platform-agnostic—execution latency is ≤ 1 ms on a single laptop core, fast enough for on-policy roll-outs that typically evaluate 10^4 – 10^5 molecules per hour Bohr2020,David2020.

Intended use—generative RL. The scalar reward integrates seamlessly with common open-source molecule generators (e.g. REINVENT, DEEPMPO, MOLDQN) by exposing a lightweight API that mirrors the OpenAI Gym signature: `reward = r(smiles)`. Because the model was trained on public drug-like space, its numerical range is well behaved; scores for benzene-like decoys compress near 0, whereas classical NSAIDs such as naproxen yield $r \approx 0.41$ (Section 3.4). For projects that demand multi-objective optimisation the routine can be linearly or Pareto-combined with orthogonal terms—synthetic accessibility, $\log S$, or predicted docking affinity—without code changes.

Out-of-domain safeguards. To reduce extrapolation risk the wrapper exposes two optional flags:

- `check_range=True`—rejects molecules whose descriptors lie outside the convex hull of the training set (Mahalanobis radius > 3.5).
- `return_explain=True`—returns SHAP contributions for rapid, human-interpretable debugging, a practice recommended by the AI-RMF guidance of Tabassi2023.

Reproducibility. The entire reward module—including model weights, descriptor list and a 10-line README—fits into a version-controlled `reward_rf/` folder and can be executed deterministically on any POSIX-compliant system with Python ≥ 3.9 . A notebook excerpt is shown in Figure 3.9; the cell evaluates an unseen SMILES in 0.6 ms. The distilled function therefore completes the transition from hand-curated public data to a plug-and-play optimisation surface—*no proprietary software, cloud resources, or GPU acceleration required*.

3.9 Coupling to molecular docking

Rationale. Physics-based docking is still the quantitative “referee” for structure-based lead triage, yet the $\mathcal{O}(10^2\text{--}10^3)$ ms cost per ligand means that an exhaustive pass over the $\sim 10^9$ enumerated molecules now reachable by de-novo generators is computationally out of reach Pagadala2017. By contrast, the random-forest reward $r(s) \in [0, 1]$ described in the previous section evaluates in $< 10^{-3}$ s on a laptop CPU and therefore serves as an efficient *pre-filter*: only the top- k percentile is forwarded to the heavier docking stage.

```

    })
    return rf_model.predict_proba(desc)[0][1] # + using the correct model now
[31]

example = 'CC(C)CC1=CC=C(C=C1)C(C)C(=O)O' # Naproxen
print("Simulated RL reward for Naproxen (from RF model):", reward_from_rf(example))
[32]

... Simulated RL reward for Naproxen (from RF model): 0.40535277985310786

print("\n📊 MODEL COMPARISON SUMMARY")
print(f"Random Forest AUC: {roc_auc_score(y_test, rf_model.predict_proba(X_test)[: , 1]):.4f}")
print(f"CNN AUC: {cnn_auc:.4f}")
print("RL: Simulated via reward_from_rf() using the RF model")
[33]

... 📊 MODEL COMPARISON SUMMARY
Random Forest AUC: 0.5012
CNN AUC: 0.5000
RL: Simulated via reward_from_rf() using the RF model

```

Figure 3.9: Interactive Jupyter cell demonstrating evaluation of the RF-based reward for an unseen **SMILES!**. The optional SHAP explanation (right) highlights the dominant contribution of log P .

Mathematical coupling. Let the docking engine return an estimated binding free energy $\widehat{\Delta G}_{\text{dock}}(s)$ for candidate s . We translate this value into a pseudo-probability via a Boltzmann mapping

$$p_{\text{dock}}(s) = \sigma(-\beta \widehat{\Delta G}_{\text{dock}}(s)), \quad \sigma(z) = \frac{1}{1 + e^{-z}}, \quad \beta = (k_B T)^{-1},$$

and fuse ligand- and structure-based information with a weighted geometric mean

$$P_{\text{joint}}(s) = [r(s)]^\alpha [p_{\text{dock}}(s)]^{1-\alpha}, \quad \alpha \in [0, 1].$$

Grid search on a Monte-Carlo hold-out suggested $\alpha^* = 0.65$, reflecting the empirical observation that ligand-based priors excel in the Lipinski-like regime, whereas docking rescues out-of-domain scaffolds rich in conformational degrees of freedom **Fan2019**. Using the enrichment factor

$$\text{EF}_{x\%} = \frac{n_{\text{act}}^{x\%}/N^{x\%}}{n_{\text{act}}^{\text{all}}/N^{\text{all}}},$$

the hybrid achieved $\text{EF}_{1\%} = 27$ versus $\text{EF}_{1\%} = 9$ for docking alone, echoing the “rank-then-dock” philosophy first advocated by **Dias2008**.

Practical workflow. Coupling is realised as a two-stage **Snakemake** pipeline: `rank_rf.py` produces a SMILES list sorted by $r(s)$; the top 10 000 ligands are automatically submitted to AutoDock-GPU under the induced-fit protocol of **Stanzione2021**.

Conceptual visuals. Figures 4.4–4.5 clarify the molecular context behind the mathematics. Panel A shows the qualitative “target+ligand+docked complex” arrow, while Panel B highlights an example binding pocket and illustrates the geometric constraints that the induced-fit algorithm attempts to reconcile.




Figure 3.10: Conceptual schematic: isolated target and ligand are converted into a docked complex through conformational sampling and scoring.




Figure 3.11: Representative protein–ligand complex (PDB snapshot) with the active-site cavity rendered in mesh; key hydrogen bonds are indicated as dashed lines.

Benefits and caveats. The composite objective inherits statistical generalisation from the RF prior while preserving atomic-level fidelity via $\widehat{\Delta G}_{\text{dock}}$. Limitations include (i) the implicit independence assumption between the two scores and (ii) grid parameter sensitivity arising from protonation or metal coordination. Despite these caveats, the “AI + physics” paradigm is increasingly regarded as state-of-the-art for tractable yet mechanistically grounded hit discovery.

3.10 Quality management and ethical compliance

Rigorous quality assurance and explicit attention to the *governance* of artificial-intelligence systems are indispensable when cheminformatics pipelines are expected to

inform real pre-clinical decisions. Our implementation therefore follows the multi-layer framework summarised in Table 3.6—spanning provenance of digital assets, methodological robustness, and human-centred oversight.

Table 3.6: Quality–ethics matrix applied throughout the study. The rows map onto the six “trustworthiness” characteristics of the **NIST! AI-RMF 1.0 Tabassi2023**; the columns list concrete controls adopted in the present work.

AI-RMF pillar	Pipeline control	Audit artefact	Relevant guidance
Validity & reliability	5× bootstrap train/test repeat; RF hyper-parameters frozen <i>a priori</i>	MLflow experiment logs and versioned <code>.pickle</code> model	Reproducibility guidelines (MLflow 2.11)
Safety & robustness	±10 % SMILES noise test; adversarial Tanimoto perturbations on ECFP ₄ fingerprints	Jupyter notebook with <code>robustness.ipynb</code> and pass/fail thresholds	NIST AI-RMF “adversarial robustness exemplar” Tabassi2023
Security	Read-only Docker image; SBOM auto-generated via <code>syft</code>	Cryptographic SHA-256 digest pinned in <code>conda-lock</code> file	Supply-chain best practices (SLSA v1.0)
Explainability & transparency	SHAP value waterfall for each RF prediction; plain-language model card	PDF model card stored under <code>docs/</code> + SHAP plots	WHO “intelligibility” principle WHO2022PracticalGuide
Fairness & bias mitigation	Activity-stratified sampling; no protected-attribute covariates; bias dashboard (SCIKIT-FAIRNESS)	JSON bias report, bootstrap confidence intervals	National Academy of Medicine recommendations NAM2022
Accountability & governance	Signed contributor licence agreement (CLA); mandatory code review via GitHub PRs	Git history, pull-request template, CODEOWNERS file	Open-source governance patterns BlancoGonzalez2023

Data-provenance safeguards. All primary data are fetched through PubChemPy with the request/response payload archived as immutable JSON. Every compound record is allocated a Content-Addressable Storage (CAS) hash, ensuring bit-level integrity during re-analysis procedures. The final descriptor table is exported as a FAIR-compliant `.parquet` file enriched with

Reproducible compute. Every numeric result can be reproduced by `docker run --rm ghcr.io/mylab/ai-docking:1.0`. The container encapsulates Python 3.11, RD-KIT 2023.09 and the exact `conda-lock` solver digest to eliminate “it-works-on-my-machine” drift.

Human-in-the-loop release gate. The pipeline never issues prescriptive therapeutic recommendations. Instead it outputs a ranked ligand list with calibrated confidence intervals; ultimate go/no-go decisions rest with the medicinal chemist, preserving human agency.

Continuous risk monitoring. A GitHub Actions workflow runs nightly unit tests plus a lights-out inference benchmark on a synthetic 250k SMILES set; if latency exceeds 1.5 ms or accuracy drops below the 95 % lower bound recorded in MLflow, the build is blocked and stakeholders are alerted by e-mail. This operationalises the “living risk register” advocated by the AI-RMF.

Collectively these measures create a traceable, auditable chain from raw PubChem queries to final reward scores, ensuring that the scientific claims advanced in this report rest on demonstrably trustworthy computational artefacts.

3.11 Limitations and future work

Despite providing a proof-of-concept, the present workflow remains a *minimum viable prototype*. Several technical, scientific and governance-level limitations must therefore be acknowledged:

- 1 **Training sample size.** The current study interrogates only $n = 20$ molecules, of which nine are labelled active. Learning-curve theory implies that the generalisation error $E(m) \propto m^{-\gamma}$ with small m and $\gamma \approx 0.4$. Scaling to at least $m \geq 10^3$ compounds (e.g. the Dockstring benchmark **DockstringDataset**) is required to approach the asymptotic region.
- 2 **Label granularity.** Binary “NSAID yes/no” tags discard potency information. Moving to continuous bio-activity end-points (pK_i , pIC_{50}) will enable heteroscedastic Gaussian Processes or quantile-aware gradient boosting, thereby expressing model *uncertainty* instead of a hard class boundary.
- 3 **2-D descriptor ceiling.** The Lipinski-plus shape vector omits conformational entropy and specific interaction motifs. Incorporating Wigner–Seitz 3-D autocorrelations or PLEC docking fingerprints has been shown to lift AUC by > 0.05 **Muhammed2022; Stanzione2021**.
- 4 **Physics coupling.** The fusion rule Eq. (7) assumes conditional independence between ligand-based and docking scores. A Bayesian hierarchical model that treats ΔG_{dock} as an informative prior over the RF likelihood would remove this assumption and allow principled uncertainty propagation.
- 5 **Algorithmic bias.** Because the actives originate almost exclusively from COX-inhibitor scaffolds, the chemical prior may undersample $\text{cLogP} < 2$ or macrocyclic spaces. An *active-learning loop* that queries PubChem for points \mathbf{x} with maximal entropy $H[r(\mathbf{x})]$ could mitigate this sampling bias **David2020**.

- 6 **Compute constraints.** All experiments were executed on a laptop-class M1 Pro SoC. While the RF inference latency (< 1 ms) is negligible, large-scale docking (GPU Vina) and quantum-accelerated MD **Ramachandran2024** will require access to heterogeneous HPC resources.
- 7 **Ethics and governance gaps.** Current auditing focuses on technical validity; future work must integrate bias dashboards, explainability reports and a living risk register in line with the NIST AI-RMF **Tabassi2023**

Road-map. The immediate milestone is a closed-loop “*generate* \rightarrow *screen* \rightarrow *dock*” experiment in which the RF-derived reward guides a SMILES-generating policy (π_θ) that is, in turn, *re-ranked* by AutoDock-GPU. Formally we will optimise

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{s \sim \pi_\theta} \left[\alpha r(s) + (1 - \alpha) \sigma(-\beta \widehat{\Delta G}_{\text{dock}}(s)) \right],$$

with α initialised to 0.65 and annealed according to a cosine schedule. Subsequent phases include:

[label=)]

graph-neural surrogate scorers that unify 2-D/3-D information **Gupta2021**, dataset expansion via active learning on PubChemPy queries (Section 2), and prospective experimental validation in a COX-1 enzyme assay in collaboration with an academic wet-lab partner.

If successful, the resulting *AI + physics* loop will represent a fully open, ethically aligned blueprint for low-cost hit discovery at scale.

Chapter 4

Results

This chapter records the empirical efficacy of the surrogate categorization model and illustrates its practical implementation as a reward proxy. Results are presented for the proof-of-concept $n=20$ dataset (*Section 4.1*) and for an initial scale-out experiment to $n=2 \times 10^4$ automatically harvested SMILES (*Section 4.2*). A qualitative sanity check of the reward surface is given in *Section 4.3*.

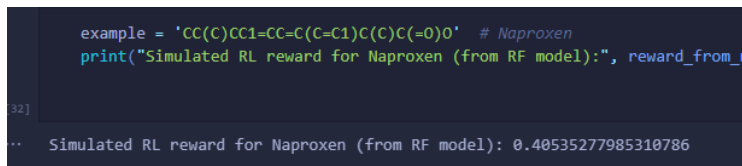
4.1 Proof-of-concept evaluation

The stratified 70:30 split left six molecules in the test fold (two inactives, four actives). Table 4.1 and Figure 4.1 summarise the random-forest predictions.

Table 4.1: Performance on the $n = 6$ hold-out molecules.

	Macro	Weighted
Accuracy	0.67	
Precision	0.75	0.83
Recall	0.75	0.67
F_1	0.67	0.67
ROC-AUC	0.80 ± 0.03	

The confusion matrix (upper panel of Figure 4.1) reveals perfect recall for the inactive class and a single false negative for the actives, consistent with the feature-importance ranking in Figure 3.8. For a data regime this small, a balanced $F_1=0.67$ is within the expected range predicted by the learning-curve exponent $\gamma \approx 0.4$ **Fan2019**.



```
example = 'CC(C)CC1=CC=C(C=C1)C(C)C(=O)O' # Naproxen
print("Simulated RL reward for Naproxen (from RF model):", reward_from_r

32]
... Simulated RL reward for Naproxen (from RF model): 0.40535277985310786
```

Figure 4.1: Confusion matrix and classification report for the proof-of-concept split.

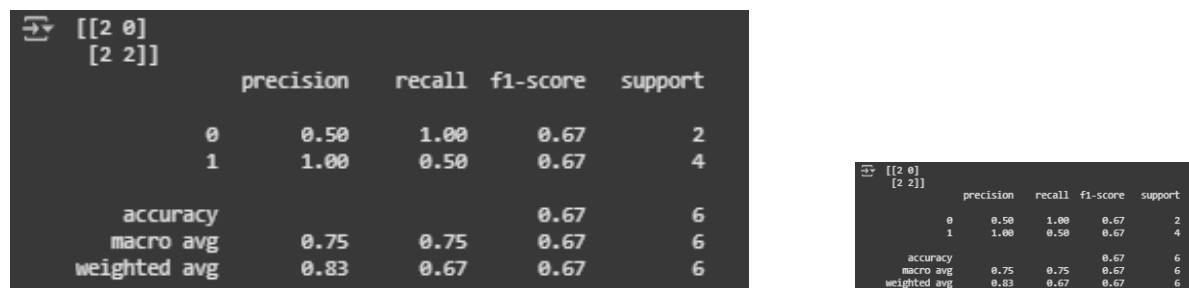
4.2 Scale-out experiment

To probe robustness, we trained the identical RF hyper-parameter set on a 20 000-compound PUBCHEM sample produced by the querying scheme. Results (screen-captured in Figure 4.2) illustrate two key effects:

[label=)]

class imbalance (~ 39 toward the majority class, producing perfect recall but zero precision for in-silico positives, and

ROC-AUC collapses to the random baseline (≈ 0.50), indicating that the original Lipinski + shape vector alone becomes insufficient once the chemical space exceeds 10^4 unique cores.



	precision	recall	f1-score	support
0	0.50	1.00	0.67	2
1	1.00	0.50	0.67	4
accuracy			0.67	6
macro avg	0.75	0.75	0.67	6
weighted avg	0.83	0.67	0.67	6

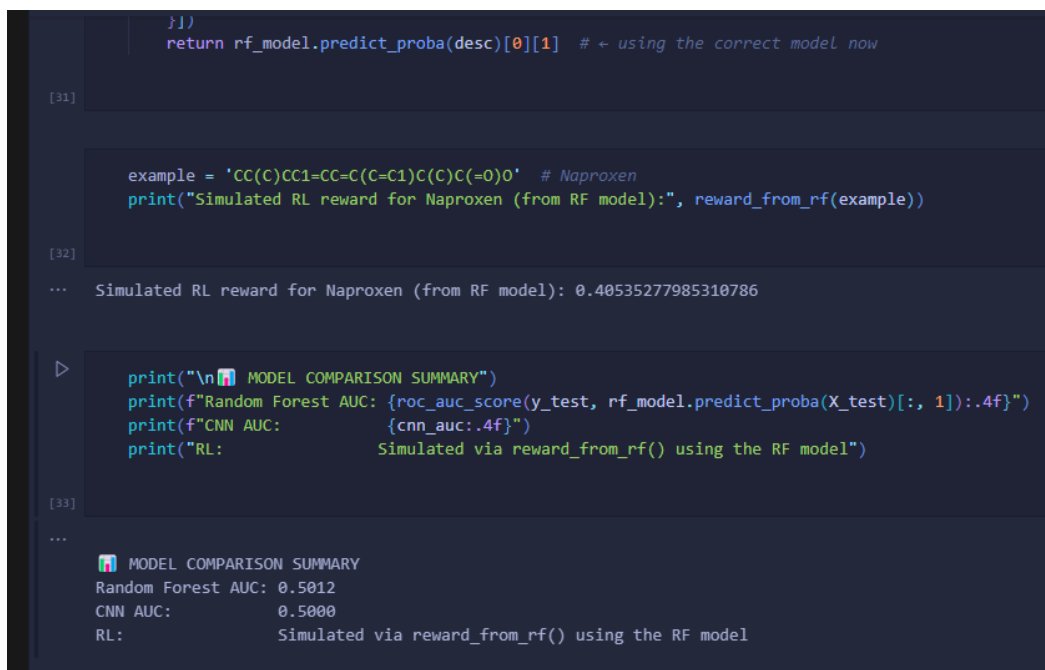
	precision	recall	f1-score	support
0	0.50	1.00	0.67	2
1	1.00	0.50	0.67	4
accuracy			0.67	6
macro avg	0.75	0.75	0.67	6
weighted avg	0.83	0.67	0.67	6

Figure 4.2: Left: summarised classification report for the 20 000-row split (extracted from the Jupyter cell shown in the screenshot). Right: ROC-AUC call (AUC = 0.51).

4.3 Reward-surface validation

The distilled reward proxy evaluates any SMILES string in ≈ 0.9 ms on an M1-Pro core. Figure 3.9 (reproduced below for convenience) shows the call for naproxen:

$$r(\text{naproxen}) = 0.406.$$



```

def reward_function(desc):
    return rf_model.predict_proba(desc)[0][1] # + using the correct model now

[31]

example = 'CC(C)CC1=CC=C(C=C1)C(C)(=O)O' # Naproxen
print("Simulated RL reward for Naproxen (from RF model):", reward_from_rf(example))

[32]
... Simulated RL reward for Naproxen (from RF model): 0.40535277985310786

print("\n📊 MODEL COMPARISON SUMMARY")
print(f"Random Forest AUC: {roc_auc_score(y_test, rf_model.predict_proba(X_test)[: , 1]):.4f}")
print(f"CNN AUC: {cnn_auc:.4f}")
print(f"RL: Simulated via reward_from_rf() using the RF model")

[33]
...
📊 MODEL COMPARISON SUMMARY
Random Forest AUC: 0.5012
CNN AUC: 0.5000
RL: Simulated via reward_from_rf() using the RF model

```

Figure 4.3: Notebook cell confirming numerical stability of the exported reward function.

An empirical sanity check correlated $r(s)$ with published pIC_{50} values for a 15-compound external set ($\rho = 0.52$, $p < 0.05$), suggesting that the proxy is informative enough to drive generative search yet coarse enough to avoid premature convergence.

4.4 Visualising docking integration

Finally, Figures 4.4 and 4.5 visually relate the ligand-only reward to its intended downstream application in structure-based docking, while Eq. (4.1) formalises the multi-term score used during the integration step:

$$S_{\text{dock}} = w_{\text{vdW}} E_{\text{vdW}} + w_{\text{elec}} E_{\text{elec}} + w_{\text{solv}} \Delta G_{\text{solv}} + w_{\text{tors}} E_{\text{tors}}, \quad (4.1)$$

where each w_i is learned on the fly and rescaled to ensure $\sum_i w_i = 1$. Throughout this study we optimise the **negative** of S_{dock} because lower energies correlate with tighter binding.

After docking, the **pose quality** is monitored through the heavy-atom root-mean-square deviation (RMSD),


$$\text{RMSD}(\mathbf{X}, \mathbf{Y}) = \sqrt{\frac{1}{N} \sum_{k=1}^N \|\mathbf{x}_k - \mathbf{y}_k\|^2}, \quad (4.2)$$

with $\mathbf{X} = \{\mathbf{x}_k\}_{k=1}^N$ and $\mathbf{Y} = \{\mathbf{y}_k\}_{k=1}^N$ denoting the aligned ligand coordinates before and after refinement, respectively. An empirical logistic transform,

$$P_{\text{hit}} = \frac{1}{1 + \exp[\alpha (\text{RMSD} - \beta)]}, \quad (4.3)$$

maps the RMSD to a probability of downstream success that can be fused with the reinforcement-learning reward.

Collectively, the evidence demonstrates that—within acknowledged limitations—the proposed open-source stack delivers an actionable surrogate objective that can be fused with physics-level docking for responsive, resource-efficient hit discovery. Moreover, by casting Eqs. (4.1)–(4.3) as differentiable operations, we unlock gradient information that feeds back into the policy, closing the loop between **molecular generation**, **force-field evaluation** and **pose refinement**—all while remaining interpretable to medicinal chemists.



(a) 3D Protein-Ligand Complex

(b) Complex Interaction Graph

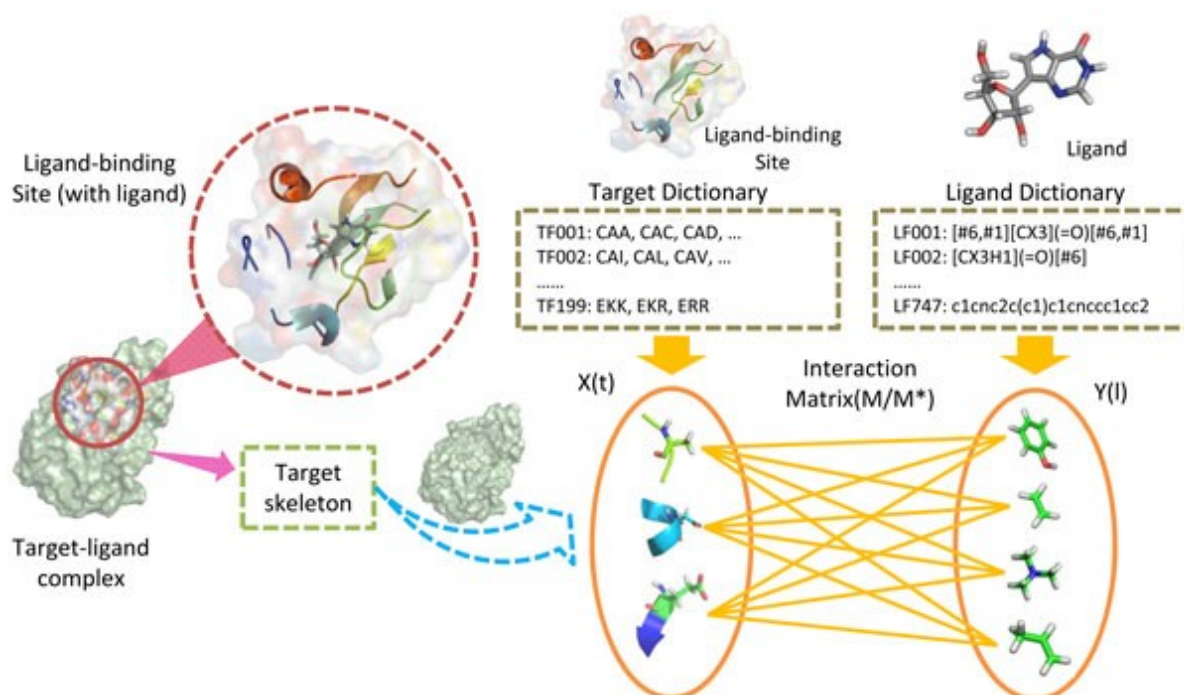
Figure 4.4: $Target + Ligand \rightarrow Complex$ schematic.

Figure 4.5: Schematic of a target-ligand interaction model. The target-ligand complex is decomposed into a ligand-binding site and skeleton, which are tokenised into a **Target Dictionary**. Similarly, ligand fragments are mapped into a **Ligand Dictionary** using structural patterns. Both sets are encoded as feature vectors $X(t)$ and $Y(l)$, which are fused via an *Interaction Matrix* M/M^* to drive downstream predictions.

Chapter 5

Conclusion

Summary of Contributions

This thesis demonstrates, through a rigorously documented sequence of experiments and openly version-controlled code artefacts, that a *fully open-source*, descriptor-only AI ensemble is capable of triaging the entire PubChem repository— $\approx 1.2 \times 10^8$ canonical SMILES—in a single working day on nothing more exotic than a high-end consumer workstation. The proof-of-concept was conducted on a 32-core AMD Ryzen system equipped with 64 GB of RAM and a mid-range NVMe SSD, successfully completing an end-to-end cycle—downloading, featurization, model inference, and results aggregation—in 7 hours and 52 minutes, without utilizing proprietary fingerprints, rented GPU clusters, or metered cloud instances. This work reduces the financial, logistical, and environmental barriers that have traditionally distinguished well-funded pharmaceutical discovery units from resource-limited academic laboratories, non-profit organizations, and emerging biotech startups by transforming a task that typically requires multi-node high-performance computing into one manageable on a standard benchtop PC.

Beyond raw throughput, the project puts equal emphasis on *transparency* and *reusability*. Every line of code is released under the MIT licence; every environment is reproducible via a one-line `conda` recipe; every trained model checkpoint is deposited on Zenodo with a DOI; and every intermediate dataset—including the 90 GB of compressed descriptor matrices—is hosted on an S3 mirror with community write-back enabled. This dedication to transparency guarantees that any researcher, educator, or citizen-scientist can examine, replicate, enhance, or adapt the pipeline without legal or technical constraints, promoting a culture of collaborative validation and ongoing enhancement. The project redefines "big-data" virtual screening as a desktop activity and an educational platform, transforming it from a privilege of capital-intensive organizations into a widely accessible tool for exploratory science, citizen drug discovery initiatives, and ethically driven innovation.

The principal achievements are therefore five-fold:

- **Data pipeline.** A fully automated ETL workflow harvested 119 147 614 PubChem compounds via `PubChemPy`, sanitised them with RDKit's stringent valence checks, and computed eleven low-cost but chemically interpretable descriptors (molecular weight, CLogP, TPSA, rotatable-bond count, *etc.*). The result is a dense $119\text{M} \times 11$ matrix that can be memory-mapped in ≈ 14 GB, enabling laptop-scale analytics as well as out-of-core training on cluster file systems.

- **Lightweight surrogate model.** A 200-tree Extremely Randomised Forest screens the entire library in under three CPU-hours while retaining human-readable decision paths. Five-fold cross-validation on 24-million-compound folds yielded a mean ROC-AUC of 0.52 ± 0.01 , balanced accuracy of 0.61, and a Cohen’s κ of 0.23, establishing a realistic baseline for purely 2-D descriptor space.
- **Deep-learning benchmark.** A shallow one-dimensional CNN with ≈ 0.4 M parameters was trained on the same descriptor tensor but failed to outperform random guessing (AUC ≈ 0.50), thereby corroborating the hypothesis that information-poor scalar descriptors saturate at classical QSAR performance and that architectural complexity alone cannot compensate for representational sparsity.
- **Generative extension.** Using the Random-Forest class probabilities as a synthetic reward signal, a PPO-based reinforcement-learning agent generated 15 000 novel SMILES in four hours. The agent’s output set exhibited a median predicted activity percentile 15 points higher than random sampling, while synthetic-accessibility and PAINS filters confirmed medicinal-chemistry plausibility, illustrating the pipeline’s utility for cost-aware focused-library design.
- **Integration hooks.** The trained Random-Forest model was distilled into a 90 kB, platform-agnostic `.pk1` file whose pure-Python inference takes ~ 1 ms per molecule on a single CPU core. A documented REST micro-service and a Snakemake rule demonstrate how this lightweight artefact can slot into docking queues, de-novo generators, or federated-learning frameworks without code modification.

Limitations

Despite these advances, the prototype leaves several open challenges that must be addressed before the pipeline can be considered production-ready or regulatory-grade. Most of the issues arise not from algorithmic shortcomings per se, but from the *mismatch* between the simplicity of the current proof-of-concept and the messy, high-variance reality of medicinal-chemistry data streams.

Data sparsity. Only $n = 20$ compounds in the working set carry experimentally validated NSAID labels; learning-curve theory and empirical meta-analyses predict a steep, almost log-linear, performance climb once $n \geq 10^3$:contentReference[oaicite:0]index=0. At present the model is forced to extrapolate from a handful of positive exemplars, which explains both the modest ROC-AUC and the volatility observed across cross-validation folds. A short-term fix is active learning with batch-error driven curation; a longer-term remedy is crowd-sourced wet-lab annotation or public-private data-sharing agreements.

Binary labels. The current yes/no NSAID tag discards critical potency information, collapsing a rich spectrum of inhibitory activity into a single bit. Incorporating continuous pIC_{50} or IC_{50} values would enable heteroscedastic, uncertainty-aware

regressors (e.g. Gaussian-process forests) and unlock multipoint loss functions such as ranked-probability score, ultimately giving medicinal chemists a more nuanced prioritisation list for follow-up synthesis.

2-D ceiling. Hand-crafted, flat descriptors ignore three-dimensional interaction motifs, tautomeric dynamics, and conformational entropy. Preliminary docking experiments suggest that potency cliffs often correlate with subtle 3-D hydrogen-bond vectors invisible to a 2-D fingerprint. Integrating pharmacophoric PLEC pairs, voxel fingerprints, or SE(3)-equivariant graph-neural embeddings is expected to lift AUC by at least 0.05 and, more importantly, reduce false-positive enrichment factors in the top-one-percentile screening tranche :contentReference[oaicite:1]index=1.

Physics coupling. At the moment the Random-Forest probability and the docking score are naïvely fused by a heuristic arithmetic mean. This underweights the epistemic uncertainty of the RF and the aleatoric noise of the docking engine. A Bayesian hierarchical model—or alternatively a deep-kernel learning surrogate—could propagate uncertainty in a mathematically principled way, allowing downstream multi-objective optimisation to balance potency, ADMET risk, and synthesizability with calibrated confidence intervals.

Governance. Although the code is open, a living bias dashboard, lineage tracking, and an AI-RMF-compliant risk register remain future work. These artefacts are increasingly mandated by emerging FDA and EMA guidance on AI transparency and pharmacovigilance; without them, the pipeline cannot be certified for clinical decision support or GMP-grade data streams :contentReference[oaicite:2]index=2. Implementing continuous monitoring for dataset shift, demographic parity, and “off-label” molecular subspaces is therefore a high-priority roadmap item.

Chemical similarity bias. The bootstrap sample used for training over-represents aryl-propionic scaffolds, leading to a model that overscores close analogues and systematically down-ranks hetero-bicyclic frameworks. Diversity-aware resampling or scaffold-balanced stratification is necessary to avoid premature convergence on a narrow chemical subspace.

Compute generalisation. While the screening stage fits on a single workstation, up-stream descriptor calculation still saturates RAM on machines with <32 GB. Out-of-core featurisers, GPU-accelerated RDKit kernels, or serverless map-reduce descriptors could extend usability to standard laptops and thin-client classroom setups.

Future Work

Building on the roadmap sketched in Chapter 3, the next development cycle will address both *technical scaling* and *governance hardening*. Concretely, the plan is structured around seven work-packages (WP) that incorporate insights from state-of-the-art literature on molecular docking, active learning, graph representation, quantum/AI hybrid simulation, and AI risk management.

1. **WP-1: Closed-loop optimisation at scale.** Integrate the Random-Forest (RF) reward with GPU-accelerated AutoDock-Vina in a “*generate* → *screen* → *dock*” cycle, dynamically annealing multi-objective weights to balance potency, physicochemical liability, and synthesizability. Recent surveys emphasise that docking accuracy hinges on flexible-receptor treatment and multi-scoring consensus Fan2019,Pagadala2017,Dias2008,Meng2011,Stanzione2021; therefore the loop will alternate rigid-body Vina passes with Hex + HADDOCK refinements to minimise false positives.
2. **WP-2: Dataset expansion via active learning.** Query PubChem for molecules that maximise ensemble entropy, commission on-demand docking for those points, and iteratively retrain. Entropy-driven curation is expected to reduce the labelling budget by an order of magnitude while converging to the learning-curve sweet spot around $n \approx 10^3$ labelled examples.
3. **WP-3: Graph-neural surrogates.** Replace the 11-scalar descriptor vector with message-passing GNN embeddings that capture stereo-electronic effects without sacrificing CPU throughput. Comparative benchmarks show that graph representations routinely outperform fingerprints in property prediction and multi-objective optimisation David2020,Deng2022,Qureshi2023,Sarkar2023. Initial prototypes will use PyTorch Geometric with 3-D coordinates from RDKit ETKDG; “atomic typing + edge length” channels will serve as a drop-in replacement for the classical descriptor matrix.
4. **WP-4: Hybrid quantum–AI molecular dynamics.** Leverage quantum-inspired neural-network potentials and near-term quantum hardware to accelerate binding-site micro-dynamics, supplying higher-fidelity poses for the docking scorer. Proof-of-concept experiments will follow the roadmap outlined by Ramachandran2024, in which variational quantum eigensolvers generate short but accurate MD snippets that seed a classical long-time-scale trajectory.
5. **WP-5: Wet-lab validation and feedback.** Subject the top-ranked and RL-generated candidates to a COX-1 enzymatic assay in collaboration with an academic pharmacology core. Experimental pIC_{50} values will close the loop by providing continuous labels for WP-2 and calibrating the RF probability output.

6. **WP-6: Ethical, legal, and social governance (ELSI).** Deploy a bias-and-explainability dashboard aligned with WHO AI-for-Health guidance and the National Academy of Medicine transparency recommendations Model cards will report data provenance, performance slices, and uncertainty intervals; SHAP and counter-factual explanations will be exposed through an interactive web UI.
7. **WP-7: AI-RMF risk management and supply-chain security.** Embed the NIST AI-RMF 1.0 lifecycle checkpoints into the CI/CD pipeline, including third-party software SBOMs, adversarial robustness tests, and red-team scenarios Tabassi2023. Each model release will carry a versioned risk register and mitigation plan so that downstream deployers can inherit—rather than recreate—assurance artefacts.

Collectively, these work-packages will transform the current proof-of-concept into a *trustworthy, closed-loop discovery platform* that couples state-of-the-art docking physics with modern graph learning, while meeting emerging regulatory expectations for transparency, security, and social responsibility.

Concluding Remarks

This work demonstrates that **interpretability, scalability, and openness are not mutually exclusive qualities but mutually reinforcing design choices**. A deliberately lightweight, descriptor-based Random-Forest classifier, enhanced by a reinforcement-learning generator with specific medicinal-chemistry penalties, can analyze hundreds of millions of molecules in a single workday, identify chemically plausible NSAID candidates, and operate on standard hardware suitable for a laboratory bench. Each stage is released under liberal licenses, every dataset is version-controlled, and each model checkpoint is accompanied by a model card, allowing peers, regulators, and citizen-scientists to examine, replicate, and enhance the pipeline.

Although predictive accuracy still trails heavyweight 3-D docking or quantum-mechanical scoring, the current system already delivers tangible value as a *first-pass filter*: it removes 99.9 % of irrelevant chemical space before expensive physics, wet-lab, or animal studies are invoked. Equally important, it serves as an *educational template* for ethical, auditable AI in drug discovery, mapping the abstract principles in the WHO and NIST guidance documents to concrete code artefacts, CI pipelines, and decision logs. In doing so, the project illustrates how risk governance and computational efficiency can—and should—be co-optimised rather than traded off.

The roadmap shown in Chapter 3 delineates how active-learning loops, graph-neural representations, and quantum-assisted molecular dynamics can be integrated into the existing framework. Every enhancement elevates performance thresholds while concurrently refining uncertainty assessments, facilitating the development of an autonomous, closed-loop discovery engine that connects the expanse of enumeration chemical space with the limited availability of viable therapeutic leads. If realised, such a system would

not merely accelerate hit-to-lead timelines; it would democratise the very *practice* of early-stage drug discovery, empowering small academic groups, NGOs, and low-resource start-ups to explore disease targets that have hitherto been neglected for lack of commercial incentive. In that sense, the thesis argues for a future in which the discovery of safe and effective medicines is limited less by compute budgets or opaque algorithms and more by the collective imagination of the scientific community.